

# Acme: A Research Framework for Distributed Reinforcement Learning

Matt Hoffman<sup>1\*</sup>, Bobak Shahriari<sup>1\*</sup>, John Aslanides<sup>1\*</sup>, Gabriel Barth-Maron<sup>1\*</sup>, Feryal Behbahani<sup>1</sup>, Tamara Norman<sup>1</sup>, Abbas Abdolmaleki<sup>1</sup>, Albin Cassirer<sup>1</sup>, Fan Yang<sup>1</sup>, Kate Baumli<sup>1</sup>, Sarah Henderson<sup>1</sup>, Alex Novikov<sup>1</sup>, Sergio Gómez Colmenarejo<sup>1</sup>, Serkan Cabi<sup>1</sup>, Caglar Gulcehre<sup>1</sup>, Tom Le Paine<sup>1</sup>, Andrew Cowie<sup>1</sup>, Ziyu Wang<sup>1</sup>, Bilal Piot<sup>1</sup>, and Nando de Freitas<sup>1</sup>

<sup>1</sup>DeepMind, \* corresponding authors; equal contribution

Deep reinforcement learning has led to many recent—and groundbreaking—advancements. However, these advances have often come at the cost of both the scale and complexity of the underlying RL algorithms. Increases in complexity have in turn made it more difficult for researchers to reproduce published RL algorithms or rapidly prototype ideas. To address this, we introduce Acme, a tool to simplify the development of novel RL algorithms that is specifically designed to enable simple agent implementations that can be run at various scales of execution. **Our aim is also to make the results of various RL algorithms developed in academia and industrial labs easier to reproduce and extend.** To this end we are releasing baseline implementations of various algorithms, created using our framework. In this work we introduce the major design decisions behind Acme and show how these are used to construct these baselines. We also experiment with these agents at different scales of both complexity and computation—including distributed versions. Ultimately, we show that the design decisions behind Acme lead to agents that can be scaled both up and down and that, for the most part, greater levels of parallelization result in agents with equivalent performance, just faster.

## 1. Introduction

Reinforcement learning (RL) provides an elegant formalization of the problem of intelligence (Russell, 2016; Sutton and Barto, 2018). In combination with advances in deep learning and computational resources, this formulation has led to dramatic results in acting from perception (Mnih et al., 2015), game playing (Silver et al., 2016), and robotics (OpenAI et al., 2018) among others. A central goal of much of this work is to create a general agent that can learn to achieve goals across a wide range of environments (Legg and Hutter, 2007). In pursuit of this objective, the scale and complexity of agents developed by the research community has grown over time; OpenAI Five (Berner et al., 2019) and AlphaStar (Vinyals et al., 2019) serve as just two recent examples of this phenomenon.

A characteristic of much of recent RL research has been an integrationist perspective on agent design, involving the combination of various algorithmic components. Agents may incorporate ideas such as intrinsic rewards and auxiliary tasks (Jaderberg et al., 2017; Jaques et al., 2019), ensembling (Osband et al., 2018), prioritized experience replay (Schaul et al., 2015), distributional backups (Bellemare et al., 2017), specialised neural network architectures (Wang et al., 2015), policy improvement search methods (Silver et al., 2018), learning from demonstrations (Hester et al., 2018; Nair et al., 2018; Gulcehre et al., 2020), variance reduction (Wang et al., 2016; Schulman et al., 2017; Espeholt et al., 2018), hierarchy (Kulkarni et al., 2016; Vezhnevets et al., 2017) or meta-learning (Al-Shedivat et al., 2017; Finn et al., 2017; Xu et al., 2018), to name a few examples. This has led to many state-of-the-art agents incorporating numerous heterogeneous components, contributing to their increased complexity and to growing concerns about the reproducibility of research (Pineau et al., 2020).

Numerous recent advances in machine learning systems have been attributable to increases in scale along two principal dimensions: function approximation capacity (number of trainable parameters) and amount of data (number and quality of training examples). In the context of RL, we focus discussion to the latter. In contrast to most supervised and unsupervised learning settings, an RL agent must interact with an environment to generate its own training data. This motivates interacting with multiple instances of an environment (simulated or otherwise) in parallel to generate more experience to learn from. This has led to the widespread use of increasingly large-scale

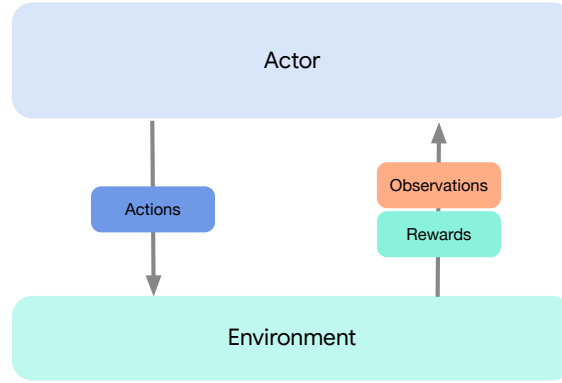


Figure 1 | A simple, high-level illustration of an actor interacting with its environment. Here we illustrate the flow of information between an actor which produces actions and the environment which consumes those actions in order to produce rewards and novel observations.

distributed systems in RL agent training (Mnih et al., 2016; Horgan et al., 2018a; Espeholt et al., 2018; Kapturowski et al., 2019). This approach introduces numerous engineering and algorithmic challenges, and relies on significant amounts of infrastructure which can impede the reproducibility of research. It also motivates agent designs that may represent dramatic departures from canonical abstractions laid out in the reinforcement learning literature (Sutton and Barto, 2018). This often means that “scaling up” from a simple, single-process prototype of an algorithm to a full distributed system may require a re-implementation of the agent.

**Acme is a software library and light-weight framework for expressing and training RL agents which attempts to address both the issues of complexity and scale within a unified framework, allowing for fast iteration of research ideas and scalable implementation of state-of-the-art agents.** Acme does this by providing tools and components for constructing agents at various levels of abstraction, from the lowest (e.g. networks, losses, policies) through to workers (actors, learners, replay buffers), and finally entire agents complete with the experimental apparatus necessary for robust measurement and evaluation, such as training loops, logging, and checkpointing. The agents written in the framework are state-of-the-art implementations that promote the use of common tools and components, hopefully leading to common community benchmarks. Our **modular** design of Acme’s agents makes them easily scalable to large distributed systems, all while maintaining clear and straightforward abstractions and **simultaneously** supporting training in the non-distributed setting.

In what remains of this section, we give a brief overview of modern reinforcement learning and discuss various software frameworks used to tackle such problems. Section 2 goes on to introduce the key structural contributions of our approach to designing RL agents. In Section 3 we build upon this design to show how this framework is used to implement a number of modern agent implementations. Finally, in Section 4 we experiment with these agents and demonstrate that they can be used to obtain state-of-the-art performance across a variety of domains.

### 1.1. Modern reinforcement learning

The standard setting for reinforcement learning consists of a learning *agent*—an entity that perceives and acts—interacting with an unknown *environment* in discrete time (Figure 1). An agent is primarily characterized by its *policy*  $\pi$ , which maps its experienced history of *observations*  $(o_0, \dots, o_t)$  to an *action*  $a_t$ . The functional form of an agent’s mapping might, for example, be represented using a feed-forward or recurrent neural network whose inputs include previous observations and actions. Once the agent has acted in the environment it then receives a reward signal  $r_t$ , makes an observation  $o_{t+1}$ , and this cycle continues. The agent’s goal is to maximize an aggregate of future rewards it expects to receive by acting upon the environment. Note that this definition of the RL problem is very broad and encompassing. For instance, the environment’s dynamics can be stochastic or deterministic, stationary or non-stationary, and include other agents. Likewise, the agent can learn from experiences generated by other behaviour policies (*off-policy*) or by its own policy (*on-policy*).

Throughout this work we will also refer to the data generation processes which interact with the environment as *actor processes* or more simply as *actors*. This is in contrast with the concept of *learners*, i.e. the processes which

consume data in order to update policy parameters, typically by stochastic gradient descent. Classically, these two processes proceed in lockstep with one another. However, by making this explicit actor/learner distinction, we can also design agents which consist either of a single actor or many distributed actors which feed data to one or more learner processes. Overall, any agent interacting within this setting has to master two formidable challenges which align with these two processes.

First, an agent must explore its environment effectively so as to obtain useful experiences. Second, it has to learn effectively from these experiences. In *online RL*, both challenges are attacked simultaneously. As a result, vast numbers of interactions are often required to learn policies represented as deep neural networks. The need for data of this magnitude motivates the use of distributed agents as described above with many parallel actors. This is particularly important in simulated environments and games where massive amounts of experience can be gathered in a distributed manner and at rates substantially faster than real-time. At the other end of the spectrum lies *offline RL*—also known as *batch RL*—which focuses on the challenge of learning policies from a fixed dataset of experiences. This situation arises often in settings where online experimentation is impossible or impractical, e.g. industrial control and healthcare. Frequently, the goal of this setting is to learn a policy that outperforms those used to generate the dataset of past experiences. Of course there is also a great deal of work in between these two extremes, which is where most of the work on off-policy agents lies.

Acme is designed to greatly simplify the construction of agents in each of these settings. In Section 2 we introduce natural modules to the design of agents which correspond to the acting, dataset, and learning components introduced above. These allow us to tackle simple, classical on- and off-policy agents by combining all of the above in a synchronous setting. We can also separate the acting and learning components and replicate the actor processes to arrive at modern, distributed agents. And by removing acting completely and making use of a fixed dataset we can tackle the offline RL setting directly. Finally, in order to exemplify this split we will also detail in Section 3 a number of example learning components built using Acme and show how these can be combined to arrive at different algorithms.


## 1.2. Related work

Numerous open-source software libraries and frameworks have been developed in recent years. In this section we give a brief review of recent examples, and situate Acme within the broader context of similar projects. OpenAI baselines (Dhariwal et al., 2017) and TF-Agents (Sergio Guadarrama, 2018) are both examples of established deep RL frameworks written in TensorFlow 1.X. They both strive to express numerous algorithms in single-process format. Dopamine (Castro et al., 2018) is a framework focusing on single-process agents in the DQN (Mnih et al., 2015) family, and various distributional variants including Rainbow (Hessel et al., 2018), and Implicit Quantile Networks (Dabney et al., 2018). Fiber (Zhi et al., 2020) and Ray (Moritz et al., 2017) are both generic tools for expressing distributed computations, similar to Launchpad, described below. ReAgent (Gauci et al., 2018) is primarily aimed at offline/batch RL from large datasets in production settings. SEED RL (Espeholt et al., 2019) is a highly scalable implementation of IMPALA (Espeholt et al., 2018) that uses batched inference on accelerators to maximize compute efficiency and throughput. Similarly, TorchBeast (Küttler et al., 2019) is another IMPALA implementation written in Torch. SURREAL (Fan et al., 2018) expresses continuous control agents in a distributed training framework. Arena (Song et al., 2019) is targeted at expressing multi-agent reinforcement learning.

The design philosophy behind Acme is to strike a balance between simplicity and that of modularity and scale. This is often a difficult target to hit—often it is much easier to lean heavily into one and neglect the other. Instead, in Acme we have designed a framework and collection of agents that can be easily modified and experimented with at small scales, or expanded to high levels of data throughput at the other end of the spectrum. While we are focusing for the moment on releasing the single-process variants of these agents, the design philosophy behind the large-scale distributed versions remains the same—as we will detail in the following section.

## 2. Acme

Acme is a library and framework for building readable, efficient, research-oriented reinforcement learning algorithms. At its core Acme is designed to enable simple descriptions of RL agents that can be run at many different scales of execution. While this usually **culminates** in running many separate (parallel) acting and learning processes in one large distributed system; we first describe Acme in a simpler, single-process setting, where acting and learning are perfectly synchronized. **A key feature of Acme** is that the agents can be run in both the single-process



<pre> 1: while True do 2:   <math>t \leftarrow 0</math>; <math>e_t \leftarrow \text{False}</math> 3:   Reset environment to get <math>o_0</math> 4:   Observe <math>o_0</math> 5:   while not <math>e_t</math> do 6:     <math>a_t \leftarrow \pi(o_t)</math> 7:     Take action <math>a_t</math>; obtain step        <math>\eta_t = (r_t, o_{t+1}, e_{t+1})</math> 8:     Observe <math>(a_t, \eta_t)</math> 9:     Update the policy <math>\pi</math> 10:    <math>t \leftarrow t + 1</math> 11:  end while 12: end while </pre>	<pre> while True:     # Make an initial observation.     step = environment.reset()     actor.observe_first(step)      while not step.last():         # Evaluate the policy and take a step in the environment.         action = actor.select_action(step.observation)         step = environment.step(action)          # Make an observation and update the actor.         actor.observe(action, next_timestep=step)         actor.update() </pre>
--	---

Figure 2 | Expansion of the environment loop to display the mechanism by which an actor interacts with the environment. Also shown for comparison is pseudocode describing this interaction as well as the actual implementation of this loop.

and highly distributed regimes using the exact same modules or building blocks with very limited differences. We achieve this by factoring the code into components that make sense at both ends of the scale. In what remains of this section we will discuss several of these components and how they interact.

## 2.1. Environments, actors, and environment loops

One of the core concepts within reinforcement learning is that of the *environment* with which an agent interacts. We will assume an environment which maintains its own state and is interacted with sequentially such that taking action  $a_t$  produces a tuple  $(r_t, o_{t+1}, e_{t+1})$  consisting of a reward, a new observation, and an end of episode indicator. Importantly, note that we have chosen to subscript each reward such that it coincides with the action that produced it. Acme assumes that the environment adheres to `the dm_env.Environment interface`. However, readers familiar with `the dm_env.TimeStep interface` may notice that we've deliberately omitted the *environmental discount factor* to simplify notation, as it often simply takes binary values to signal the end of an episode.

In Acme, the component that interacts most closely with the environment is the *actor*. At a high level, an actor consumes observations produced by the environment and produces actions that are in turn fed into the environment. Then, after observing the ensuing transition, we give the actor an opportunity to update its internal state; this most often relates to its action-selection policy, but we will return to this point shortly.

`The interaction between an actor and its environment is mediated by an environment loop`. Custom loops can easily be implemented but we provide a generic one that meets most of our needs and provides a simple entry point for interacting with any of the actors or agents implemented within Acme. In Figure 2 we illustrate this interaction in further detail by expanding the loop (shown earlier in Figure 1) to include the exact interface by which an actor interacts with its environment. Given an observation  $o_t$  we must first be able to evaluate the

actor’s action-selection policy  $a_t = \pi(o_t)$ , where  $a_t$  can also represent a sample of some random variable in the case of a stochastic policy. Once an action is taken the actor must be able to record the reward and subsequent observation obtained from the environment—e.g. one might insert this data into a replay table or collect an entire trajectory to be processed at the end of an episode. These two methods are included in our illustration, and show the life-cycle of a single iteration of an environment loop. This figure also shows the pseudocode and includes a (slightly simplified) example of Acme’s implementation of this loop. We stress here, however, that while this formulation is not a new concept—it can be found in any introductory text (e.g. Sutton and Barto, 2018)—it does serve to highlight one of the key design goals of Acme: **wherever possible there should be a one-to-one mapping between typical RL pseudocode and its implementation.**

## 2.2. Learners and Agents

Thus far we’ve focused our attention on components that are relevant for data generation: acting in the environment and observing ensuing transitions. We now introduce the **learner** component, which consumes this data in order to obtain a better policy. This component often contains the bulk of the code relevant to any specific RL algorithm and, in deep RL, takes the form of optimizing the weights of a neural network to minimize some algorithm-specific loss(es). More precise mathematical descriptions for a variety of algorithms will be detailed in Section 3. While it is possible to run a learner without further interaction with the environment (see e.g. Section 2.6), in RL we are often interested in concurrent learning and acting. Therefore **we introduce a special type of actor that includes both an acting and a learning component; we refer to these as *agents* to distinguish** them from their non-learning counterparts.

While an agent defers its action selection to its own acting component, its update method, elided from the previous figure, is where an agent triggers some number of learning steps within its learner component. In contrast, a generic actor’s update method simply pulls neural network weights from a variable source if it is provided one at initialization. Since a learner component is a valid variable source, the actor component may query a learner directly for its latest network weights. This will be particularly relevant when we discuss distributed agents in Section 2.4.

In Figure 3 we again show the environment loop, where we have expanded the interaction to show the internals of a learning agent we just described. While **redundant**, we sometimes use the term *learning agent* to emphasize that the agent contains a learner component. The illustration includes the actor and learner components and depicts how they interact. In particular, the actor pulls weights from the learner components in order to keep its action-selection up-to-date. Meanwhile, the learner pulls experiences observed by the actor through a dataset, which is yet another important component.

## 2.3. Datasets and Adders

Note that having a dataset component sit between the actor and learner components is quite a general view that includes on- and off-policy learning, and experience replay—prioritized or otherwise—depending on how the dataset is configured. From the learner’s perspective data is provided simply as a stream of sampled mini-batches; the dataset can be configured to hold on to stale data, and/or the actors can be programmed to add noise to the learner-specified policy. **While we have generally standardized on TensorFlow’s `Dataset` object to provide efficient buffering and iteration over data,** this does not mandate the use of TensorFlow for the update step implemented by a learner. The dataset itself is backed by a low-level data storage system, called Reverb (Albin Cassirer, Gabriel Barth-Maron, Manuel Kroiss, Eugene Brevdo, 2020), that is released concurrently. Reverb can be roughly described as a storage system which enables efficient insertion and routing of items and a flexible sampling mechanism that allows: first-in-first-out, last-in-first-out, uniform, and weighted sampling schemes.

Acme also provides **a simple common interface for insertion into the low-level storage system in the form of *adders*.** Adders provide add methods which are analogous to the observe methods found on an actor—in fact most actors’ observations are forwarded directly onto an adder object. These objects exist in order to allow for different styles of pre-processing and aggregation of observational data that occurs before insertion into the dataset. For example a given agent implementation might rely on sampling transitions,  $n$ -step transitions, sequences (overlapping or not), or entire episodes—all of which have existing adder implementations in Acme.

Relying on different adder implementations to carry the workload of adding data once it is observed has also allowed us to design very general actor modules that support a wide variety of learning agents. While any agent is

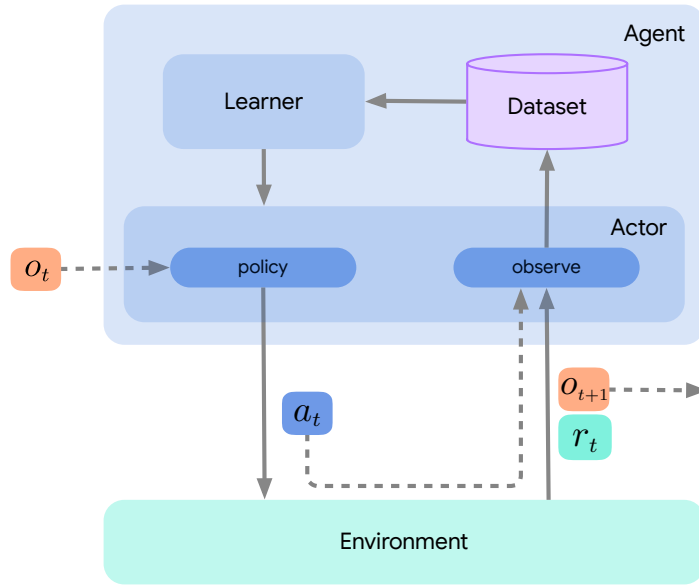


Figure 3 | Example of the environment loop with a full learning agent.

able to implement its own internal actor—or indeed bypass its actor component entirely and implement its own acting/observing methods directly as an agent is an actor in its own right—most agents are able to use one of these standard actors. Actors in Acme predominantly fall into one of two styles: feed-forward and recurrent. As their names suggest, these actors primarily differ in how they maintain state (or do not) between calls to the action-selection method, and the exact form of network used for these actors must be passed in at construction. Note that in Acme we have also taken pains to ensure that the communication between different components is agnostic to the underlying framework (e.g. TensorFlow) used. However, as the actors themselves must interact directly with this framework we also provide different implementations for both TensorFlow and JAX—and similar accommodations could be made for other frameworks.

Given these different modules we can easily construct novel algorithms by varying one or more components. For example, the *distributional discrete policy gradient* method we will introduce later consists of a feed-forward actor, an  $n$ -step transition adder, the distributional DPG losses, and either uniform or prioritized sampling under the dataset. However, easily composing modules in order to create novel agents is not the primary purpose of these components. Instead, in the next section we will describe how these modules can easily be pulled apart at the boundaries in order to enable distributed agents that can run at much larger scales.

## 2.4. Distributed agents

Up to this point we have primarily described the interaction between an actor—or agent—and its environment using a simple, synchronous setting. However, a common use case is to generate data asynchronously from the learning process, often by interacting with multiple environments in parallel (Nair et al., 2015; Mnih et al., 2016; Horgan et al., 2018a; Barth-Maron et al., 2018; Kapturowski et al., 2019). In Acme we accomplish this by splitting the acting, learning, and storage components introduced earlier into different threads or processes. This has two benefits: the first being that environment interactions can occur asynchronously with the learning process, i.e. we allow the learning process to proceed as quickly as possible regardless of the speed of data gathering. The other benefit gained by structuring an agent in this way is that by making use of more actors in parallel we can accelerate the data generation process.

An example of a distributed agent is shown in Figure 4. By examining this figure we can see that this largely maintains the same structure introduced earlier. However, in the previous section links between different modules were used merely to indicate function calls. Instead, in the distributed variant each module is launched in its own process, where the links between different modules are now used to illustrate *remote procedure calls* (RPC). In the illustrated example this distributed agent consists of a data storage process, a learner process, and one or more






Figure 4 | Example of a distributed, asynchronous agent. In contrast to Figure 1 we have moved the the replay and update components into external processes here designated by grey nodes. We have also replicated the environment loop and replaced the actor itself by a thin *proxy actor* which pulls parameter updates from a learner.

distributed actors, each with their own environment loop. In order to simplify this construction, we also frequently use an additional sub-module on the actor process: a *variable client*. This serves purely to allow the actor to poll the learner for variable updates and simplifies the code (which is more cumbersome if the learner is required to push to every actor).

In this work we focus on describing the single-process variants of Acme agents. As a result, a full description of the distributed agents is somewhat out of scope. However, here we briefly describe the tool we have developed, **Launchpad**, which enables these distributed variants. The agents introduced in the previous section were themselves composed of different sub-modules, e.g. an actor, learner, and data storage system. In the same way, Launchpad can be thought of as a mechanism for composing these modules in the distributed setting. Roughly speaking, Launchpad provides a mechanism **for creating a distributed program as a graph consisting of nodes and edges**. Nodes exactly correspond to the modules—represented as class instances as described above—whereas the edges represent a client/server channel allowing communication between two modules. Once this graph has been constructed the program can then be launched to start its underlying computation. The key innovation of Launchpad is that it handles the creation of these edges in such a way that from **the perspective of any module there is no distinction between a local and remote communication**, e.g. for an actor retrieving parameters from a learner in both instances this just looks like a method call.

As a result, in what remains of this work we will primarily describe either the individual modules, or their single-process combinations. Our results, however, will show both single-process and distributed variants and in both cases the same underlying learning and acting code is being used. We will leave a further detailed description of Launchpad for later work.

## 2.5. Reverb and rate limitation

First introduced by (Lin, 1992), **experience replay** has since been successfully applied to deep reinforcement learning (Mnih et al., 2013), **allowing agents to reuse previous experiences**. In Acme observations are added to the replay buffer through the **actor's `observe()` method**, which is called with each time step. Batches of transitions used to train the agent are then sampled from the buffer using **the learner's `update()` method**. Designing a replay buffer requires careful consideration regarding how to package experience into elementary items, how to sample

these items, and how to remove them when the buffer is full. With Reverb, these features are easily configured allowing the agent code to focus on what behaviour to use rather than how to achieve it.

In a synchronous learning loop, you may prescribe how many steps of acting in the environment your agent should perform between each learning steps. This ratio between acting and learning has a dramatic effect on not only the sample efficiency (the number of environment steps required to reach a given performance) but also the long-term learning performance and stability. The same is true for distributed learning settings although this setting makes it more difficult to maintain a fixed ratio. Indeed if distributing an agent is for the sake of computational efficiency, then it is clearly not desirable to block the learning process while the actor processes are gathering data. On the other hand, running both processes independently easily results in higher variance. The variance is often attributable to differences in the computational substrate (e.g. different hardware and network connectivity) between seeds but pinpointing precise sources can be extremely challenging.

In Acme, these scaling issues are mitigated through the use of Reverb’s `RateLimiter`. By adopting rate limitation, one can enforce a desired relative rate of learning to acting, allowing the actor and learner processes to run unblocked so long as they remain within some defined tolerance of the prescribed rate. In an ideal setting, both processes are given enough computational resource to run unblocked by the rate limiter. However if due to network issues, insufficient resources, or otherwise, one of the processes starts lagging behind the other, the rate limiter will block the latter while the former catches up. While this may waste computational resource by keeping the latter idle, it does so only for as long as is necessary to ensure the relative rate of learning to acting stays within tolerance.

Indeed the replay buffer is a good place to enforce such a rate. Notice that the first step in a learning process is to sample data from the dataset; similarly a key step in the environment loop (which runs on the actor processes) is to observe transitions and insert data into the dataset. Since both the learning and acting processes must communicate with the dataset component. If the learner is sampling experiences too quickly, the buffer blocks sampling requests until the actor catches up; if the actor(s) are inserting experiences too quickly, the buffer blocks insert requests until the learner catches up.

## 2.6. Offline Agents

An additional benefit of the structuring we have taken for Acme agents is that it is trivial to apply any learning component to the offline setting, which assumes a fixed dataset of experiences that can be learned in a purely supervised manner. In Acme this is as simple as applying a given learner module to a given dataset of experiences. As there is a great deal of overlap between the pure batch setting and off-policy learning, many off-policy agents perform quite well and/or can be adapted to work when not allowed to interact with the environment. As we will describe in Section 3, however, agents purpose-built for this setting can have the edge when considering the underlying distribution of generated data. We defer any further discussion of this use-case to Section 3.7.

## 3. Agent Implementations

In this section, we describe a set of agents whose implementations we include in the initial release of Acme. Our intention is for these agents to serve both as clear and succinct *reference implementations* of their respective RL algorithms, as well as strong research *baselines* in their own right. More agents may become available in the future but we hope that these serve as readable examples for the RL community to base their own design and implementations on. We will also relate these agents back to the underlying components of Acme used to construct these agents. However, for most of this section the primary differences between agents will be contained within the *learner*, i.e. the module which consumes data and updates the agent’s policy parameters.

In what follows, we will describe the underlying learning algorithms for the agents that we are releasing along with Acme. While this is not intended as a full tutorial, we will first provide a brief introductory background in order to keep this section reasonably self-contained. We will then roughly organize our discussion around the following groupings of agents: temporal difference learning and value-based agents, followed by policy optimization and actor-critic agents.



### 3.1. Background and notation

The interaction between an *agent* and its *environment* can typically be formalized as a (partially-observable) Markov decision process ((PO)MDP Cassandra et al., 1994; Puterman, 1994). This formalism corresponds to the *environment loop* introduced earlier (see Figure 2). However, for clarity we repeat our earlier discussion, i.e. at time  $t$  the agent makes *observation*  $o_t$  given by the environment and selects action  $a_t = \pi(o_t)$  as determined by its policy. Upon taking this action the environment emits *reward*  $r_t$  and transitions to a new state, resulting in observation  $o_{t+1}$ . This interaction happens indefinitely or until the environment produces an episode *termination signal*  $e_{t+1}$ .

For simplicity we will also denote both deterministic and stochastic policies with  $\pi$  and we will simplify such policies to depend only on the most recent observation. To fully solve this problem in general, partially-observable domains the policy may be required to depend on the entire history of observations  $o_{0:t}$ . To address this, modern implementations frequently take the approach of using a learned summary statistic, i.e. a recurrent state, that is also output by the policy network. Some algorithms we introduce below will make use of this more complicated formulation, however for the sake of simplicity we will only do so when necessary.

While each algorithm we introduce below consumes data generated by an environment loop as above, one way in which they differ is in the format that data takes as it is presented to the learner. The simplest—and perhaps most classical—form in which data is exposed is as a *transition*  $(o_t, a_t, r_t, o_{t+1})$ . By collecting an entire sequence of transitions until the termination signal  $e_t$  is true we can also form episodes  $(o_t, a_t, r_t)_{t \geq 0}$  or sub-slices of the episode which we refer to as *sequences*. As noted above these individual elements will be exposed to the learning algorithm using the interface of a dataset, typically by sampling at random (uniformly or with some probability). However, we will also see that it is possible to process elements in the order in which they are observed, in which case the dataset takes the form of a queue. In Acme each of these data formats is handled by a relevant *adder* object.

The objective of all the agents we will cover is to maximize some form of its expected *return*. Although other aggregation methods are possible, we will focus on the sum of discounted, future rewards

$$R_t = \sum_{i \geq 0} \gamma^i r_{t+i} = r_t + \gamma R_{t+1} \quad (1)$$

where  $R_t$  is a random variable that depends on the future trajectory of the agent. While the policy  $\pi$  is probably the most important component of any RL algorithm, for all of the agents we will discuss, the *state-action value* or *Q-function* is almost as important. In the first set of algorithms we will examine this will provide either a direct parameterization of the policy, or will be indirectly used to optimize the policy. At its heart, this function, in conjunction with the policy  $\pi$ , maps any observation/action pair to the future rewards of taking that action and then following the given policy. This can be written as

$$\begin{aligned} Q^\pi(o_t, a_t) &= \mathbb{E}[R_t \mid o_t, a_t], \\ &= r_t + \gamma \mathbb{E}_\pi[Q^\pi(o_{t+1}, a_{t+1}) \mid o_t, a_t] \end{aligned} \quad (2)$$

where by returning to the recursive definition of the returns given in (1) we arrive at the celebrated *Bellman equation*. This definition allows us to start with an arbitrary approximation to  $Q^\pi$  and repeatedly, recursively improve upon that estimate. Under certain regularity conditions (see Sutton and Barto (2018) for more details) we can combine this update with updates to the policy in order to obtain the optimal *Q-function* and hence the optimal policy. In the next section we will first discuss algorithms for which the value function and the underlying policy are one-and-the-same.

### 3.2. Policy evaluation and value based agents.

While exactly computing  $Q^\pi$  using the Bellman equation is not possible when the underlying environment dynamics are unknown (and must be sampled) we can instead empirically approximate this function using observed data. This is commonly addressed in deep RL by using a neural network  $Q_\phi$  to approximate  $Q^\pi$ . In order to optimize the parameters of this network, it is common to introduce a *bootstrap target*:

$$y_{\pi, \phi'} = r_t + \gamma \mathbb{E}[Q_{\phi'}(o_{t+1}, \pi(o_{t+1}))]. \quad (4)$$

This roughly corresponds to one step of the backup operation introduced in the recursive Bellman equation. As is common practice since (Mnih et al., 2015), these targets use an identical network—dubbed the *target network*—with “stale” parameters  $\phi'$ . The *online network* (so-called in order to distinguish it from the target network) has parameters  $\phi$  and can be fit to this bootstrap target by minimizing the squared *temporal difference* (TD) error

$$L(\phi) = \mathbb{E} [y_{\pi, \phi'} - Q_{\phi}(o, a)]^2. \quad (5)$$

The expectation above is taken with respect to transitions  $(o, a, r, o')$  generated by the agent’s policy; in practice this typically means that the loss is formed empirically from samples taken from a replay buffer filled by the actor process(es). Note that as the learner trains the value function  $Q_{\phi}$ , it periodically communicates these weights to the actor. We next consider value-based agents, whose greedy deterministic policies are directly derived from the Q-function.

**Deep Q-Networks.** The first algorithm we consider is that of Deep Q-Networks (DQN) (Mnih et al., 2013, 2015). Equipped with the Q-function and facing an observation  $o$ , a greedy actor can simply select the action that maximizes its value. This indirectly defines the policy as  $\pi(o) = \arg \max_a Q(o, a)$ , where typically actions  $a$  are restricted to a finite, integer-valued set. By plugging this policy into Equation (5) we arrive at the loss used by DQN, where in this case the bootstrap target is simply a function of the sampled transition and the target parameters  $\phi'$  (this is a point that will become more important shortly). In order to optimize this loss, DQN fills a replay buffer with transitions generated in an  $\epsilon$ -greedy manner (i.e. with probability  $\epsilon$  we generate actions purely at random). This buffer is then sampled uniformly at random to form a minibatch of sample transitions which are then used to perform stochastic gradient descent on the given loss.

Note that DQN is an off-policy algorithm, meaning that is capable of learning from data generated *off-policy*, i.e. from a policy separate from the one it is optimizing. The is to distinguish from on-policy algorithms for which data generation uses the policy being optimized. While most of the algorithms that we will consider are off-policy, we will note variations from this norm as necessary.

In our implementation of DQN, and following in the spirit of Rainbow DQN (Hessel et al., 2018), we also include a number of recent enhancements. The first of these is the use of Double Q-learning (van Hasselt et al., 2015) to combat over-estimation of the Q-function. While subtle, this corresponds to a modification of the bootstrap target to

$$y_{\phi, \phi'} = r_t + \gamma Q_{\phi'}(o_t, \pi_{\phi}(o_t)). \quad (6)$$

Note that here we have also been careful with the subscripted parameters in use by each component. We see that the policy selecting the action uses the online network weights  $\phi$  while the network used to evaluate said action uses the target weights  $\phi'$ . Importantly, we do not allow gradients to flow through the bootstrap target, and as a result it is necessary to employ a *stop gradient* since the target now depends on  $\phi$ . This was not necessary with the simpler DQN variant where the target only depended on  $\phi'$ .

Our implementation also makes use of  $n$ -step targets to fit the agent’s value function—this allows the algorithm to use longer sequences of the observed reward signal when bootstrapping. In order to enable this the actor must store overlapping  $n$ -step transitions of the form  $(o_t, a_t, \sum_{i < n} \gamma^i r_{t+i}, o_{t+n})$  rather than a single transition. In this setting the target becomes

$$y_{\phi, \phi'} = \sum_i \gamma^i r_{t+i} + \gamma^n Q_{\phi'}(o_{t+n}, \pi_{\phi}(o_{t+n})), \quad (7)$$

and in Acme there is an adder built expressly for this purpose. The use of  $n$ -step transitions are also useful computationally because they are functionally equivalent to single-step transitions—which we can see by letting  $n$  equal one—and use the same amount of storage. This is a standard improvement we will see in use by many algorithms whose original variants were based on transitions. Additional enhancements to our implementation include duelling networks (Wang et al., 2015) and prioritized experience replay (Schaul et al., 2015), wherein priorities are used to sample transitions from replay proportional to their TD error rather than uniformly. Finally, the distributed variant, which is not yet open-sourced, resembles that of Ape-X DQN (Horgan et al., 2018b); with each actor running its own  $\epsilon$ -greedy behavior policy, with  $\epsilon$  drawn from a log-uniform distribution.

**Recurrent DQN.** The recurrent replay distributed DQN (R2D2) algorithm (Kapturowski et al., 2019) further extends the work of Ape-X DQN by making use of a recurrent network—in particular by incorporating an additional LSTM layer in the  $Q$ -network. This results in a value function of the form  $Q_\phi(o, a, s)$  which additionally takes a recurrent state  $s$  which must also be initialized at the beginning of an episode. The change to a recurrent network leads to several other modifications to the  $Q$ -learning process. First, rather than learning from transitions R2D2 instead relies on full  $n$ -step sequences of the form  $(o_{t:t+n}, a_{t:t+n+1}, r_{t:t+n-1})$  and using strided, double  $Q$ -learning over these fixed length sequences. In Acme this is accomplished simply by making use of a sequence adder as opposed to the transition adder used by DQN. Additionally, these sequences are sampled from replay with priorities by a convex combination of their mean and maximum absolute TD-errors. R2D2 also makes use of a transformed loss, introduced in (Pohlen et al., 2018), rather than using clipped rewards.

Finally, in order to learn the value function, R2D2 also requires a sequence of recurrent states  $s_{t:t+n}$ . This can prove problematic for sequences which do not begin from the initial state. In order to solve this problem R2D2 stores old sequence in the replay buffer alongside the sequences of observations, but with this solution comes another problem: stale recurrent states, i.e. states which are different from those currently being generated by the network. This problem is solved by storing such sequences with an extra “burn-in” period at the beginning where no learning is done but which is only used to initialize the recurrent state.

Note that while we will also present results on the distributed variant of this algorithm, the version released is not distributed. However, for simplicity we will still refer to it as R2D2—a convention we will maintain for other algorithms introduced later which were originally published under a “distributed” moniker.

### 3.3. Policy optimization and actor-critic agents

The agents we have discussed thus far have relied upon the value function  $Q_\phi$  to indirectly parameterize their policies, i.e. by selecting the value-maximizing action for a given observation. We will now discuss methods for which the policy itself is directly parameterized by weights  $\theta$  and denoted  $\pi_\theta$ . In doing so we will introduce the so-called actor-critic learning paradigm wherein a learned value function (hereafter referred to as the *critic*) is learned in tandem with the policy and is used to define the policy’s loss.

Typically, actor-critic methods rely on the same TD error minimization introduced in the previous section to optimize the critic. However, while different agents optimize different policy losses, but they are generally derived from a common goal, namely to maximize expected return

$$J(\theta) = \mathbb{E}_o \mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(o, a)] \quad (8)$$

given the value-function  $Q^{\pi_\theta}$  associated with this policy. By making use of the *policy gradient theorem* (Sutton et al., 2000) an unbiased estimate of the policy gradient can be written as

$$\nabla J(\theta) = \mathbb{E}_o \mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(o, a) \nabla \log \pi_\theta(a | o)], \quad (9)$$

where  $o, a$  and their ensuing return are sampled from the experiences of  $\pi_\theta$ . Though unbiased, these gradients can exhibit high variance. A common strategy to attack this variance is by introducing and subtracting a *baseline* independent of the action.

The most common baseline is the state value function  $V^{\pi_\theta}(o) = \mathbb{E}_a Q^{\pi_\theta}(o, a)$ . Unlike the state-action value,  $V^{\pi_\theta}$  estimates the expected return starting from the observation  $o$  and acting according to  $\pi$  thereafter—it integrates over all actions according to  $\pi_\theta$  rather than allowing for an initial deviation. The difference between these two value estimates, i.e.  $Q^{\pi_\theta}(o, a) - V^{\pi_\theta}(o)$ , is known as the *advantage* and as a result by introducing this baseline we arrive at a family of advantage actor-critic algorithms (Mnih et al., 2016; Espeholt et al., 2018) using the following policy gradient

$$\nabla J(\theta) = \mathbb{E}_o \mathbb{E}_{a \sim \pi_\theta} [(Q^{\pi_\theta}(o, a) - V^{\pi_\theta}(o)) \nabla \log \pi_\theta(a | o)]. \quad (10)$$

Algorithms in this family primarily differ in how they estimate this advantage, see Schulman et al. (2016) for a few examples.

**Importance Weighted Actor-Learner Architecture (IMPALA).** IMPALA (Espeholt et al., 2018) is a distributed advantage actor-critic agent which makes use of the loss introduced above. In order to do so it estimates the

state-action value function using Monte Carlo rollouts and produces an estimate  $V_\phi$  of the state value function by minimizing a variant of the TD error. As a result, IMPALA is very close to an on-policy algorithm and is able to make use of very long sequences. The catch is that IMPALA is designed to work in the distributed setting where the behavior policy generating this data may not exactly match the policy being evaluated.

In order to counteract this off-policy bias, a typical approach is to employ an importance sampling correction. Given a trajectory generated by a behaviour policy  $\pi_b$ , the importance sampling ratio is the ratio between the probability of the action under target policy  $\pi_\theta$  and the behaviour policy:  $\rho_t = \frac{\pi_\theta(a_t|o_t)}{\pi_b(a_t|o_t)}$ . This ratio allows us to adjust the probabilities for any generated action, reweighted as if they were actually sampled according to the target policy  $\pi_\theta$ . Although this allows us to have an unbiased estimate of the gradients, it suffers from high variance. To compensate for this, IMPALA introduces a *V-trace* correction which further truncates the importance sampling weights. This has the effect of *cutting* a trajectory once it becomes too far off-policy. This results in a recursive definition of the Bellman update at step  $t$  which can be rewritten for the state value function as

$$v_t = V_\phi(o_t) + \sum_{k=t}^{t+n-1} \gamma^{k-t} \prod_{i=t}^{k-1} c_i \rho_k (r_k + \gamma V(o_{k+1}) - V_\phi(o_k)). \quad (11)$$

We can see that the final term is simply the TD error for step  $k$  and  $c_t = \min(1, \rho_t)$  is the *truncated* importance sampling weight. As with DQN, the critic for IMPALA is updated to minimize the squared TD error, i.e. the difference between the current estimate  $V_\phi$  and the updated version  $v_t$ ,

$$L(\phi) = \mathbb{E} [v_t - V_\phi(o_t)]^2. \quad (12)$$

To update the policy IMPALA makes use of the policy gradient under this advantage estimate can then be updated to follow the entropy-regularized policy gradient. For a single step  $t$  this gradient corresponds to the first term of the following:

$$\nabla J(\theta) \approx \mathbb{E} \left[ \log \pi_\theta(a_t|o_t) (r_t + \gamma v_{t+1} - V_\phi(o_t)) - \beta \sum_a \pi_\theta(a|o_t) \log \pi_\theta(a|o_t) \right]. \quad (13)$$

The additional *entropy regularization* term is introduced to prevent instability caused by the policy moving too quickly and is a common technique used to stabilize policy gradient methods. Off-policy correction and slowing down the policy for stability are the two key ingredients of existing powerful actor-critic agents, including Actor-Critic with Experience Replay (ACER) by Wang et al. (2016), Trust Region Policy Optimization (TRPO) by Schulman et al. (2015) and Proximal Policy Optimization (PPO) by Schulman et al. (2017).

### 3.4. Continuous control

We now turn to a collection of agents where the actions taken by the agent are real valued and therefore continuous. The set of agents we consider are also actor-critic methods, but in this case making use of the state-action value function. While this is not strictly necessary, in this setting taking an arg max over values is often not tractable so actor-critic learning an attractive alternative. The following learning algorithms alternate gradient steps optimizing the critic and policy losses. As with the DQN agent, these maintain target parameters  $\phi'$  and  $\theta'$  which are periodically copied from their online counterparts for training stability. Meanwhile one or more actors will fill the replay buffer from which samples are taken in order to optimize the associated loss function. Departing from the DQN work introduced earlier, these actors make use of exploration noise of given by a Gaussian of fixed scale (i.e. rather than an epsilon-greedy strategy).

**Deep Deterministic Policy Gradient (DDPG).** Unlike IMPALA where the policy is stochastic with discrete actions, DDPG (Lillicrap et al., 2016) uses a deterministic policy with continuous actions. Our implementation employs the standard  $n$ -step TD loss in Equation (5) to train the critic  $Q_\phi$ . However, because of the deterministic policy, Silver et al. (2014) derived a new policy gradient theorem resulting in the following

$$\nabla J(\theta) = \mathbb{E}_o \left[ \nabla_\theta \pi_\theta \nabla_a Q_\phi(o, a) \Big|_{a=\pi_\theta(o)} \right]. \quad (14)$$

As noted above, our DDPG implementation roughly follows the same strategy that of DQN given above, with a replaced learner and Gaussian exploration noise. I.e. the mechanism by which data is added is equivalent. Our implementation uses uniform sampling from replay, as we have found that prioritization provides minimal (if any) benefit. We will follow this same strategy for the remaining continuous control algorithms for which the only changes necessary come to their learning strategy and losses.

**Maximum a posteriori Policy Optimization (MPO).** Introduced by Abdolmaleki et al. (2018), the MPO algorithm takes a two-pronged approach to policy optimization, inspired by the classical expectation-maximization (EM) algorithm. Because of this particular approach, its expected return takes the following peculiar form:

$$J_{\eta, \alpha}(\theta) = \mathbb{E}_o \mathbb{E}_{a \sim \pi_{\theta'}} \left[ \exp \left( \frac{Q_{\phi'}(o, a)}{\eta} \right) \log \pi_{\theta}(a|o) \right] + \alpha [\epsilon - \mathbb{E}_o D_{\text{KL}}(\pi_{\theta'}(\cdot|o) \parallel \pi_{\theta}(a|o))], \quad (15)$$

which includes a Kullback-Leibler (KL) divergence regularization that targets a hyperparameter  $\epsilon$ . This regularization makes sure (i) that the online policy  $\pi_{\theta}$  does not move too far from the target network  $\pi_{\theta'}$  and (ii) that the online policy keeps adapting if necessary. Finally, the dual variables  $\eta$  and  $\alpha$  are not hyperparameters, they are *dual variables* with losses of their own to change them adaptively.

**Distributional critics, D4PG, and DMPO.** Introduced by Bellemare et al. (2017), the C51 agent’s critic  $Z_{\phi}$  estimates the *distribution* over returns, in contrast to the critics described so far, which only estimate the *expected* return. Therefore, naturally we should have  $Q_{\phi}(o, a) = \mathbb{E} Z_{\phi}(o, a)$ . The D4PG (Barth-Maroon et al., 2018) and our novel DMPO agent are adaptations of DDPG and MPO, respectively, to use this distributional critic. Their policy losses remains almost unchanged—the critic simply needs to be averaged. For instance, the D4PG policy gradient becomes:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_o \left[ \nabla_{\theta} \pi_{\theta} \nabla_a \mathbb{E} Q_{\phi}(o, a) \Big|_{a=\pi_{\theta}(o)} \right]. \quad (16)$$

Meanwhile the critic loss has to be changed from (5) to take into account the fact that we are now working with distributions.

$$L(\phi) = \mathbb{E}_{\tau} H(Y_{\phi'}(\tau), Z_{\phi}(s, a)), \quad (17)$$

$$Y_{\phi'}(\tau) = \Pi \left[ r + \gamma Z_{\phi'}(o', \pi_{\theta'}(o')) \right], \quad (18)$$

where  $H(\cdot, \cdot)$  represents the cross-entropy. The projection  $\Pi$  is needed since the bootstrap target is now a distribution denoted by a capital  $Y_{\phi'}$  and it must be projected onto the fixed support of the critic  $Z_{\phi}$  before the cross-entropy can be computed.

### 3.5. Monte Carlo Tree Search

This agent implements planning with an environment model (learned or otherwise), with search guided by policy and value networks. This can be thought of as a scaled-down and simplified version of the AlphaZero algorithm (Silver et al., 2018), optionally coupled with an environment transition model, similar to the set-up in (van Hasselt et al., 2019). We learn a value function  $V_{\phi}$  via TD-learning and policy  $\pi_{\theta}$  via imitation of a model-based MCTS policy  $\pi_{\text{MCTS}}$ , whose search is guided by the model-free policy  $\pi_{\theta}$ :

$$a_{\text{search}} = \arg \max_a \left[ Q(o_t, a) + \beta \frac{\sqrt{N(o_t)}}{N(o_t, a) + 1} \pi_{\theta}(a | o_t) \right], \quad (19)$$

where  $N$  is a simple visit-count for a given node in a tree, and  $\beta$  is the UCT hyperparameter controlling how greedy the search should be with respect to action-value estimates. When doing tree search roll-outs, we truncate at some fixed search depth and bootstrap from our model-free value estimate  $V_{\phi}$ . The agent’s final policy is softmax with respect to value over the children of the root node of the search tree, and the model-free policy loss is given by

$$J(\theta) = D_{\text{KL}}(\pi_{\theta}(\cdot|o_t) \parallel \pi_{\text{MCTS}}(\cdot|o_t)). \quad (20)$$

### 3.6. Learning from demonstrations

When a task requires a long sequence of correct actions before obtaining any reward from the environment, it is said to be a hard exploration task. Exploration in such sparse reward tasks is particularly challenging due to the problems of reward discovery and credit assignment. Promising recent attempts at overcoming exploration in such tasks leverage demonstrations of successful trajectories (Kim et al., 2013; Hester et al., 2017; Vecerik et al., 2017; Nair et al., 2018; Gulcehre et al., 2020), dubbed the RL with Expert Demonstrations (RLED; Piot et al., 2014). In the RLED framework, an expert teacher (e.g. a human or a scripted agent) demonstrates the task using the same action space. The resulting trajectories are stored in an expert replay buffer, which is then interleaved with the learning agent’s own experiences in its minibatches. These agents must therefore necessarily learn off-policy from a replay buffer, which in practice—and in our case—is often prioritized (Schaul et al., 2015). Applying the RLED approach to DQN and R2D2, produces the **DQfD** and **R2D3** agents, respectively (Hester et al., 2017; Gulcehre et al., 2020).

### 3.7. Offline learning

In situations where interactions with the environment are expensive or dangerous, such as robotics, self-driving cars, and health-care, it is not possible to directly use RL algorithms, which gather data while they learn—classically known as *online* learning. As a result, there has been increasing interest in methods for learning policies from logged data known as offline or batch RL (Lagoudakis and Parr, 2003; Lange et al., 2012) methods. Recently, these methods have produced promising results in simple domains and there is ongoing research efforts to scale these algorithms to more challenging problems (Fujimoto et al., 2018; Agarwal et al., 2019; Cabi et al., 2019).

As mentioned earlier, one of the advantages of the modularity of Acme agents is that any of our agents can be used in the offline setting by simply providing our *learner* components a fixed dataset of experience. Indeed given a *dataset* the learning algorithm is independent of the data generation process, making it amenable to both online and offline learning. In this case, there are no actors in the distributed system and only the dataset and learner remain. However, it may be helpful to keep an evaluator process (an actor that doesn’t add it experience to replay) to continually quantify the performance of the learning algorithm. This does not necessarily mean that a learning process designed for online, or even off-policy, data will perform well in the purely offline setting—and often, we see that methods purposely-built to account for the distribution of data perform better.

One notable learning algorithm that serves as a baseline algorithm for all offline algorithms is known as **Behaviour Cloning** (BC; Pomerleau, 1989; Michie, 1993). Indeed this is the simplest form of imitation learning where the agent learns to mimic the demonstrations by learning a mapping between observations and actions via supervised learning. While it can be quite competitive when the dataset is large and includes high-quality demonstrations, many interesting applications do not have that luxury.

## 4. Experiments

Here, we present the performance of several Acme agents on various simulated environments. We put a significant effort in implementing readable, modular, and scalable agents; we present these results to demonstrate that this has not come at the expense of performance. Indeed our agents achieve returns that are comparable to the published state-of-the-art. In relating these results, we are not necessarily interested in comparing agents as much as comparing performance across distributed scales. In addition to performance plots we also study the effect of rate limitation on both sample efficiency and walltime. Before diving into the results of our benchmarks, let us briefly introduce the relevant environments and measurements that we report on.

### 4.1. Environments

**DeepMind Control suite.** The DeepMind Control Suite (Tassa et al., 2018) provides a set of continuous control tasks in MuJoCo (Todorov et al., 2012) and has been widely used as a benchmark to assess performance of continuous control algorithms. The tasks vary from the simple control problems with a single degree of freedom (DOF) such as the cartpole and pendulum, to the control of complex multi-joint bodies such as the humanoid (21 DOF). We will be considering two variants of these tasks: learning from raw features or from pixels. When learning from raw features the observations are scalars representing positions and velocities; these vary in size between 3 to 137 dimensions depending on the task. When learning from pixels, the observations are stacks of 3

consecutive RGB images of size  $72 \times 96$ , stacked along the channel axis to allow the agents to learn to use dynamic information like velocity and acceleration. We clustered these tasks into 4 categories according to the complexity of control (trivial, easy, medium and hard) running each for a sufficient number of actor steps.

**Arcade Learning Environment.** The Arcade Learning Environment (ALE) (Bellemare et al., 2013) provides a simulator for Atari 2600 games. ALE is one of the most commonly used benchmark environments for RL research. The action space ranges between 4 to 18 discrete actions (joystick controls) depending on the game. The observation space consists of  $210 \times 160$  RGB image. We use a representative subset of ten Atari games to assess the performance of our discrete agents. We also use several pre-processing methods on the Atari frames including giving zero-discount on life loss, action repeats with frame pooling, greyscaling, rescaling to  $84 \times 84$ , reward clipping, and observation stacking following (Mnih et al., 2015).

**Behaviour Suite.** The Behaviour Suite for Reinforcement Learning (bsuite) (Osband et al., 2020) provides a collection of experiments that investigate the core capabilities of RL agents. These tasks are intended to be challenging but simple and interpretable tests of specific axes of capability for a given RL agent. Each experiment studies the scalability and robustness of agents in various types of tasks, such as exploration, memory, generalization, and robustness to noise. This provides an powerful experimental apparatus for testing hypotheses about the capabilities of various RL algorithms and agents. For example we will clearly see later that those agents with recurrence (e.g. R2D2) are the only agents able to properly perform in tasks that require memory. While this conjecture can readily made from an algorithmic perspective, bsuite allows us to actually quantify the amount of gains brought by such an algorithmic component.

## 4.2. Measurements

All our agents (including the single-process ones, for fair walltime comparison) are equipped with a background evaluator node that periodically queries the learner for policy network weights and evaluates them by running an episode in the environment and logging the observed episode return. For all tasks in the control suite, an episode corresponds to 1,000 steps and the reward signal is such that 100 is the theoretical limit on the episode return. There is no such common episode return scale across Atari levels but we do limit the max episode length on the evaluator to 108,000 environment steps as is common.

**Sample efficiency and actor steps.** In many applications, sample efficiency is a critical obstacle in the way of an RL algorithm’s viability. For example stepping the environment could involve a computationally expensive physics simulator or a real robot which is slow and endures some wear-and-tear. In our possibly distributed agents, the most accurate measure of environment interactions during training is what we refer to as *actor steps*, to distinguish them from the evaluator steps. In the Atari environment, since we repeat actions 4 times, as is now standard, our measurement of actor steps is exactly four times smaller than the number of environment frames, which is a more common measure in the ALE literature.

**Speed and learner walltime.** Of course in some applications the environment may be a cheap simulator, of which any number of copies can be quickly spun up and driven in parallel by distributed actors. Though sample efficiency may still come into play, in these cases one is often more interested in how quickly results can be obtained by increasing the number of distributed actors. Unfortunately simply measuring walltime via timestamp is a rather noisy way of keeping time as it can lead to vastly different results due to uncontrollable factors. Indeed distributed computation on a shared cloud service can bring its own challenges in the form of interrupted processes and/or communication, but in the following benchmarks we are interested in the performance of distributed agents on dedicated hardware. In order to simulate dedicated hardware, we accumulate time within the learner from immediately after the first learner step; we refer to this measure of time as the *learner walltime*. Since the learner is checkpointed along with the networks, this timekeeping persists through interruptions on the learner.

## 4.3. Control suite from features

We begin by reporting our results on the control suite from features. For practical reasons, in this section we highlight a challenging subset of tasks, but a comprehensive set of benchmarks can be found in the appendix.






Figure 5 | Performance of control agents on the control suite from raw features. Comparing the single process agents. Depicted curves represent rolling averages over 10 seeds. Performance as measured by episode return with respect to actor steps.

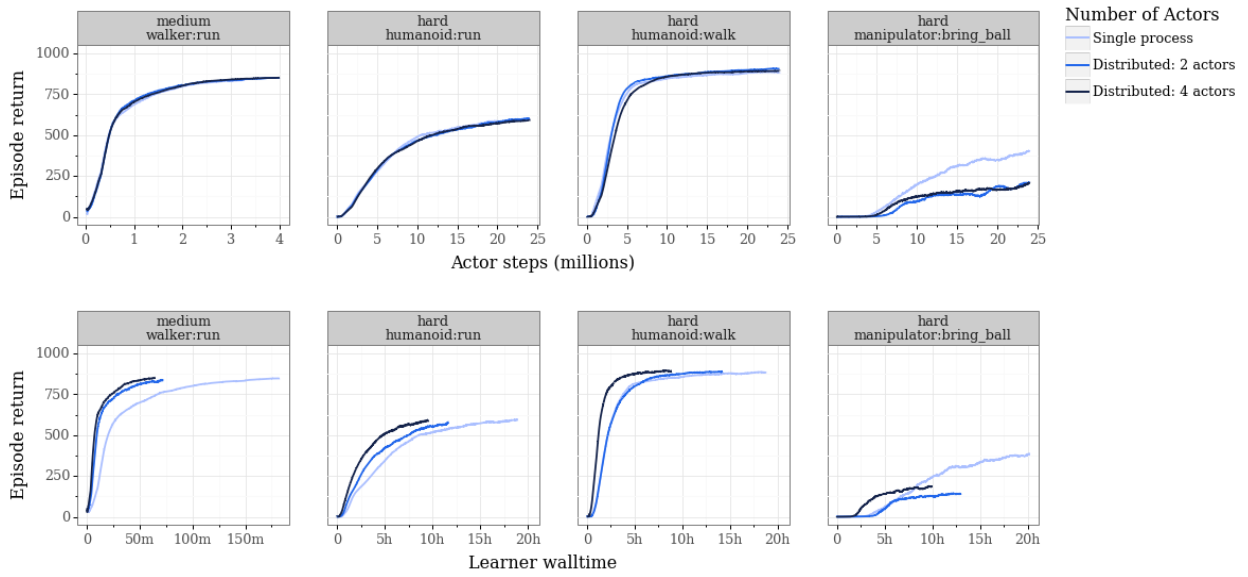


Figure 6 | Performance of D4PG on the control suite from raw features. Comparing the single-process and two distributed variants (2, 4 actors). Depicted curves represent rolling averages over 10 seeds. *Top*: Sample efficiency as measured by episode return with respect to actor steps. *Bottom*: Learning speed as measured by episode return with respect to learner walltime.

Figure 5 shows that our released single-process agents achieve state-of-the-art performance. Note that we have made reasonable efforts to tune these agents so that their defaults are suited for solving control suite tasks from features.

Let us focus on one of the agents from Figure 5, say D4PG, and let us overlay the performance of the exact same code running in a distributed setup with 2 and 4 actors, resulting in Figure 6. The top half of the figure shows that, with the exception of `manipulator:bring_ball` (a very challenging task), all variants of the agent achieve nearly identical results when measured against actor steps. Indeed this is the rate limiter’s function. Meanwhile, the bottom half of the figure shows the exact same training curves plotted against learner walltime; these show the benefit of distributed acting. Together these form compelling evidence that rate limitation is a tremendous tool for reproducibility and the fair comparison of agents between disparate computational distribution strategies.

#### 4.4. Rate limiter sensitivity analysis

Next we compare the performance of our D4PG agent while varying the rate of learning to acting. (Note that a similar analysis could easily be carried out for any of our agents.) The rate is measured in number of samples out of replay per insert into the replay buffer: samples per insert ratio (SPI). For reference, the default D4PG agent




Figure 7 | Sensitivity analysis of the effect of rate limitation on the sample efficiency of D4PG. All curves correspond to D4PG with 4 distributed actors and represent rolling averages over 10 seeds. See Figure 18 in the appendix for results on more tasks.

uses a batch size of 256 and a SPI of 32; this corresponds to doing one gradient update every 8 acting steps. These results strongly suggest that setting this relative rate of learning to acting can dramatically affect sample efficiency and performance.

Figure 7 shows a representative subset of our experiment. Overall we notice a trend that lower SPI ratios lead to more wasteful agents in the sense that they need many more interactions with the environment to reach the same performance. However we can also differentiate three distinct effects on the results depending on the task. On the `acrobot:swingup` task, the higher the ratio the better—more learning per acting leads to a more sample efficient agent. Meanwhile on the humanoid tasks, this is not the case. Indeed, we see that our chosen default of 32 SPI seems to be the most efficient choice in the long run. Though the 128 SPI agent learns fastest to begin with, it seems to have settled on a sub-optimal solution; whereas the rate of 512 SPI is clearly unstable.

Finally, notice that with no rate limitation at all, performance is similar to our default of 32 SPI. This is due to our choice of number of distributed actors (4) in this experiment, which proved to be an efficient use of computational resources.

We would like to investigate this phenomenon further in the future, and hopefully thanks to the agents and tools we are releasing, the community can help. Our current hypothesis is the following. In the `acrobot` tasks, it is possible to accurately evaluate the value of a state-action pair independent of the policy used thereafter; therefore with a reasonable number of transitions, the task can be effectively learned *offline* and certainly off-policy. Meanwhile, in the humanoid tasks, due to the complex composition of joints and the temporally extended nature of the tasks, accurately evaluating a state-action pair (which is the critic’s function) is tightly coupled with the policy being evaluated. For instance, in `humanoid:run`, moving a humanoid’s ankle to let its body fall forward is only good if a hip flexor will later bring the opposite leg forward to support the falling body.

**Off-policy experience and the role of replay buffer size.** What complicates matters further is the role of the replay buffer capacity: the larger the buffer, the more stale the experience it contains. Therefore, combining a large replay buffer with a large SPI ratio, leads to a very off-policy learning. At the opposite end of the two spectra, by using a relatively small replay buffer and SPI ratio, learning can be entirely on-policy. In order to simplify this analysis, in all these experiments, we kept the replay buffer fixed at a large size of 1 million items (in this case an item is an  $n$ -step transition).

#### 4.5. Control from pixels

We shift our attention to learning control suite tasks from pixel observations. In this setting we only ran D4PG and only on the trivial, easy, and medium tasks. Remarkably, by simply adding a residual network (ResNet) torso to the exact same network architecture as above, many complex tasks were learned without *any* additional tuning. Naturally, there is still a lot of room for improvement as many of the tasks are not learned. Here we highlight a few tasks on which D4PG performed very well; the full set of results can be found in Figures 19–20 in the appendix. Once again we see that, when measured with respect to actor steps, performance and sample-efficiency is very similar across the variants of the agent: single-process and distributed with 2, 8 and 16 actors. However, when measured with respect to the learner’s walltime, the advantage of having more actors is clear. In fact the benefits are more dramatic than in Figure 6, which is a consequence of the added computational load on each individual

actor, due to the environment rendering each frame and the additional ResNet forward pass required to process them. This experiment shows that parallelism helps accelerate training, but that the same performance can still be achieved with fewer actors and more time.

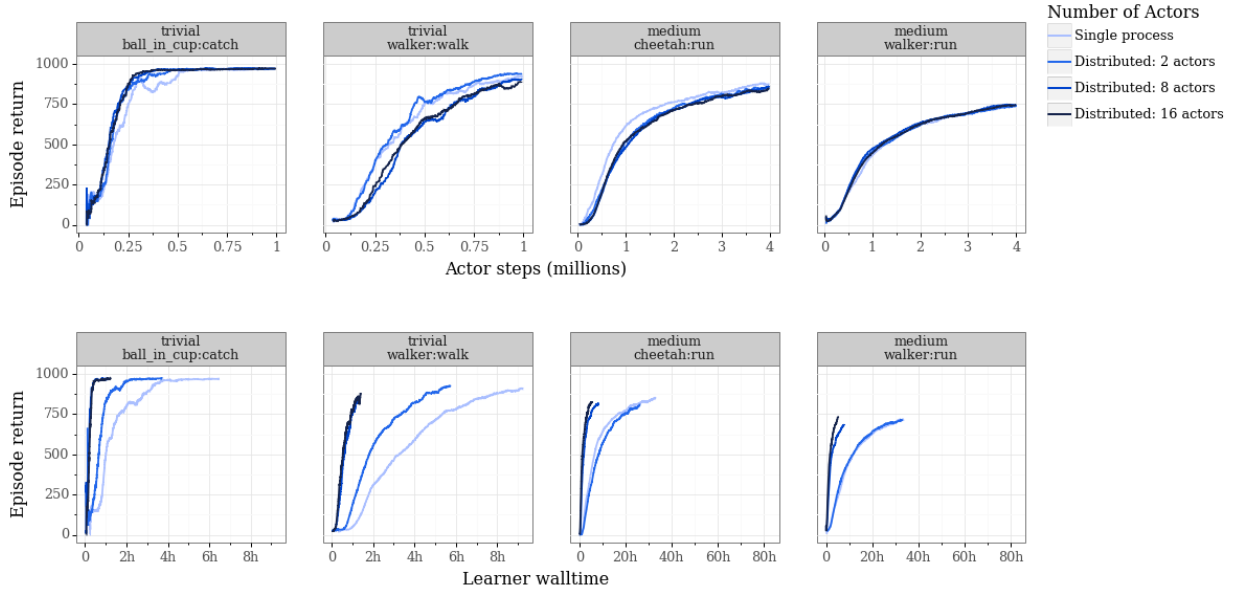


Figure 8 | Performance of D4PG on the control suite from pixel observations. Comparing the single-process and three distributed variants (2, 8, 16 actors). Depicted curves represent rolling averages over 10 seeds. *Top*: Sample efficiency as measured by episode return with respect to actor steps. *Bottom*: Learning speed as measure by episode return with respect to learner walltime.

#### 4.6. Atari

We evaluate the performance of DQN, R2D2 and IMPALA agents trained on 10 individual levels with varying levels of difficulty over 3 seeds. We use 256 actors in all cases. We have done agent-specific hyper-parameter tuning. For each agent the hyper-parameters are the same across all games and set as default parameters in Acme.




Figure 9 | Comparison of performance (sample complexity) of DQN, R2D2 and IMPALA on a subset of Atari tasks. Here we plot rolling averages across 3 seeds as measured by episode return with respect to actor steps. Note the qualitatively different learning traits for each algorithm: IMPALA typically learns quickly but is prone to instabilities; in contrast R2D2 learns slowly but typically attains higher final performance with less variance; finally DQN, being a feed-forward agent, tends to get “off the ground” faster than R2D2 but performance plateaus at a much lower level.

We report performance measured by episode return with respect to actor steps in Figure 9. Each actor step results in 4 environment steps here as we set the action repeat to 4 for all agents. All agents are trained for approximately same duration. Note that in IMPALA we are using the Deep architecture proposed in (Espeholt et al., 2018) which is more costly at runtime and hence performs fewer updates. We also measure episode return with respect to training time which can be found in Figure 21 in the appendix.

#### 4.7. bsuite

We evaluate the performance of DQN, R2D2, IMPALA and MCTS agents on the bsuite benchmark. We aggregate the performance of these agents across all tasks shown in the “radar plot” in Figure 10. We did not tune the hyperparameters and believe these results can be improved with careful tuning. In particular, IMPALA performs poorly in bandit-like domains without extra tuning due to instabilities arising from short sequences. In these experiments, MCTS has access to a perfect simulator for all tasks.




Figure 10 | Comparing aggregate performance of DQN, R2D2, MCTS and IMPALA on bsuite. Note (a) in this experiment MCTS has access to a perfect simulator; (b) IMPALA performs poorly in bandit-like domains without extra hyperparameter tuning, due to instabilities arising from short sequences.

#### 4.8. Learning from demonstrations




Figure 11 | Comparing aggregate performance of DQN, DQfD, R2D2 and R2D3 on bsuite. DQfD and R2D3 have access to demonstrations for the exploration tasks, so this comparison is only meant to quantify how well these algorithms can leverage that additional information.

To verify our DQfD and R2D3 agents, we test them on bsuite (Figure 11). We only include demonstrations for the exploration tasks. These demonstrations were generated using the optimal policy, which has knowledge of the action mapping of the environment. See code for details. For Deep Sea, one demonstration is sufficient. For Deep Sea Stochastic, because the stochastic nature of the environment, we need more demonstrations. We

generate  $\text{num\_demos} = \text{environment size} * 10$ . Because the optimal policy does not always solve the task in this environment, we include 80% successful trajectories, and 20% unsuccessful trajectories.

#### 4.9. Offline RL

We provide results on control from features and Atari offline-RL datasets. The details of those datasets will be available in an upcoming offline RL benchmark dataset release as *RL Unplugged*. RL Unplugged will include a diverse set of challenging Acme compatible offline RL datasets along with a proper evaluation method for each dataset. In this paper, we focus on the easier version of the control dataset and only a small subset of Atari games that will be released with RL Unplugged. Here, we are only interested in showing that it is possible to run offline RL agents with the Acme infrastructure. RL Unplugged will dive into a more in-depth analysis of these and other tasks.

**Control from features.** In this section we briefly show the results of BC and D4PG algorithms used in the offline setting on two control suite tasks: *cartpole:swingup* and *fish:swim* from features information (similar to Sec. 4.3). To generate the dataset, we ran D4PG in online mode with three random seeds till convergence (we call the resulting policy the *data generation policy*), collected all the data experienced by these three runs (which includes low quality data e.g. from the beginning of training), and subsampled this data, leaving 200 transitions for the *cartpole:swingup* environment and 8000 transitions for the *fish:swim* environment. Then, this data was assembled into a dataset and used for training BC and D4PG agents from scratch in offline mode, i.e. without having access to the environment during training (but still using the environment for evaluation). D4PG was able to almost match the performance of the data generation policy on both domains, but in the *cartpole:swingup* experiment its performance decreased after a while, probably due to overfitting to a small dataset (Fig. 12).

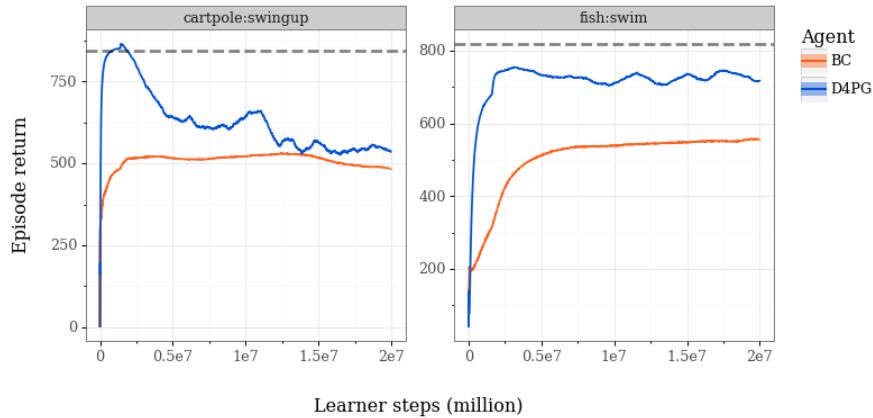


Figure 12 | Comparing performance of BC and D4PG on an offline setting with a fixed dataset as measured by episode return with respect to learner steps. The performance of the data generation policy is shown with the dashed horizontal line.

**Atari.** Atari has been an established benchmark for offline-RL (Agarwal et al., 2019; Fujimoto et al., 2019). Here, we report results on nine atari games to show that the performance of our offline Acme DQN agent can match the the best behavior policy that generated the data while being trained online. We have run our DQN atari agent on the data generated the dataset. We have used the atari offline RL dataset generated in (Agarwal et al., 2019), and trained a Double DQN (van Hasselt et al., 2015) with Adam optimizer (Kingma and Ba, 2014). We trained our agents on nine Atari games that can be categorized in terms of diverse range of difficulty which is determined in terms of the performance of the online agent that generated the datasets: RoadRunner, IceHockey, Zaxxon, DemonAttack as the easy games, BeamRider, MsPacman, Robotank, Pooyan as the medium difficulty and DoubleDunk as a hard game. We trained our agents for 5 million learner steps with minibatches of size 256 and report the results in Figure 13. Our results closely matches the results reported in (Agarwal et al., 2019) for Offline DQN on most games with the same network architecture.




Figure 13 | Comparing performance of DQN on an offline setting with a fixed dataset as measured by episode return with respect to learner steps. The best performance of the online policy during that generated the data over the course of training is shown with the dashed horizontal line.

## 5. Conclusion

In this work we have introduced Acme, a modular light-weight framework that supports scalable and fast iteration of research ideas in RL. Acme naturally supports both single-actor and distributed training paradigms and provides a variety of agent baselines with state-of-the-art performance. In this release we are focusing on the single-actor setting which is more in line with the standard needs of the academic community. However, this work also describes the design of Acme which enables the use of agent components in both these settings. Although we provide results for both settings, we also show that the same results can be obtained using single-process agents enabled by our framework.

By providing these tools, we hope that Acme will help improve the status of reproducibility in RL, and empower the academic research community with simple building blocks to create new RL agents. Additionally, our baselines should provide additional yardsticks to measure progress in the field. We are excited to share Acme with the research community and look forward to contributions from everyone, as well as feedback to keep improving and extending Acme.

## Acknowledgements

We’d like to thank Yannis Assael and Vaggelis Margaritis for a great deal of graphical feedback on this work. We’d also like to thank Jackie Kay, David Budden, and Siqi Liu for help on an earlier version of this codebase.

## References

- Abdolmaleki, A., Springenberg, J. T., Tassa, Y., Munos, R., Heess, N., and Riedmiller, M. (2018). Maximum a posteriori policy optimisation. *arXiv preprint arXiv:1806.06920*.
- Agarwal, R., Schuurmans, D., and Norouzi, M. (2019). An optimistic perspective on offline reinforcement learning. *NeurIPS Deep Reinforcement Learning Workshop*.
- Al-Shedivat, M., Bansal, T., Burda, Y., Sutskever, I., Mordatch, I., and Abbeel, P. (2017). Continuous adaptation via meta-learning in nonstationary and competitive environments. *arXiv preprint arXiv:1710.03641*.
- Albin Cassirer, Gabriel Barth-Maron, Manuel Kroiss, Eugene Brevdo (2020). Reverb: An efficient data storage and transport system for ml research. <https://github.com/deepmind/reverb>. [Online; accessed 01-June-2020].
- Barth-Maron, G., Hoffman, M. W., Budden, D., Dabney, W., Horgan, D., TB, D., Muldal, A., Heess, N., and Lillicrap, T. P. (2018). Distributed distributional deterministic policy gradients. In *6th International Conference on Learning*



- Representations, ICLR 2018, Vancouver, BC, Canada, April 30 - May 3, 2018, Conference Track Proceedings.* OpenReview.net.
- Bellemare, M. G., Dabney, W., and Munos, R. (2017). A distributional perspective on reinforcement learning. In Precup, D. and Teh, Y. W., editors, *Proceedings of the 34th International Conference on Machine Learning, ICML 2017, Sydney, NSW, Australia, 6-11 August 2017*, volume 70 of *Proceedings of Machine Learning Research*, pages 449–458. PMLR.
- Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. *J. Artif. Intell. Res.*, 47:253–279.
- Berner, C., Brockman, G., Chan, B., Cheung, V., Debiak, P., Dennison, C., Farhi, D., Fischer, Q., Hashme, S., Hesse, C., Józefowicz, R., Gray, S., Olsson, C., Pachocki, J., Petrov, M., de Oliveira Pinto, H. P., Raiman, J., Salimans, T., Schlatter, J., Schneider, J., Sidor, S., Sutskever, I., Tang, J., Wolski, F., and Zhang, S. (2019). Dota 2 with large scale deep reinforcement learning. *CoRR*, abs/1912.06680.
- Cabi, S., Gómez Colmenarejo, S., Novikov, A., Konyushkova, K., Reed, S., Jeong, R., Zolna, K., Aytar, Y., Budden, D., Vecerik, M., Sushkov, O., Barker, D., Scholz, J., Denil, M., de Freitas, N., and Wang, Z. (2019). Scaling data-driven robotics with reward sketching and batch reinforcement learning. In *Robotics, Science and Systems*.
- Cassandra, A. R., Kaelbling, L. P., and Littman, M. L. (1994). Acting optimally in partially observable stochastic domains. In *AAAI*, volume 94, pages 1023–1028.
- Castro, P. S., Moitra, S., Gelada, C., Kumar, S., and Bellemare, M. G. (2018). Dopamine: A Research Framework for Deep Reinforcement Learning.
- Dabney, W., Ostrovski, G., Silver, D., and Munos, R. (2018). Implicit quantile networks for distributional reinforcement learning. In Dy, J. G. and Krause, A., editors, *Proceedings of the 35th International Conference on Machine Learning, ICML 2018, Stockholmsmässan, Stockholm, Sweden, July 10-15, 2018*, volume 80 of *Proceedings of Machine Learning Research*, pages 1104–1113. PMLR.
- Dhariwal, P., Hesse, C., Klimov, O., Nichol, A., Plappert, M., Radford, A., Schulman, J., Sidor, S., Wu, Y., and Zhokhov, P. (2017). OpenAI Baselines. <https://github.com/openai/baselines>.
- Espeholt, L., Marinier, R., Stanczyk, P., Wang, K., and Michalski, M. (2019). Seed rl: Scalable and efficient deep-rl with accelerated central inference.
- Espeholt, L., Soyer, H., Munos, R., Simonyan, K., Mnih, V., Ward, T., Doron, Y., Firoiu, V., Harley, T., Dunning, I., Legg, S., and Kavukcuoglu, K. (2018). IMPALA: scalable distributed deep-rl with importance weighted actor-learner architectures. In Dy, J. G. and Krause, A., editors, *Proceedings of the 35th International Conference on Machine Learning, ICML 2018, Stockholmsmässan, Stockholm, Sweden, July 10-15, 2018*, volume 80 of *Proceedings of Machine Learning Research*, pages 1406–1415. PMLR.
- Fan, L., Zhu, Y., Zhu, J., Liu, Z., Zeng, O., Gupta, A., Creus-Costa, J., Savarese, S., and Fei-Fei, L. (2018). Surreal: Open-source reinforcement learning framework and robot manipulation benchmark. In *Conference on Robot Learning*.
- Finn, C., Abbeel, P., and Levine, S. (2017). Model-agnostic meta-learning for fast adaptation of deep networks. In *International Conference on Machine Learning*, pages 1126—1135.
- Fujimoto, S., Conti, E., Ghavamzadeh, M., and Pineau, J. (2019). Benchmarking batch deep reinforcement learning algorithms. *arXiv preprint arXiv:1910.01708*.
- Fujimoto, S., Meger, D., and Precup, D. (2018). Off-policy deep reinforcement learning without exploration. *arXiv preprint arXiv:1812.02900*.
- Gauci, J., Conti, E., Liang, Y., Virochsiri, K., Chen, Z., He, Y., Kaden, Z., Narayanan, V., and Ye, X. (2018). Horizon: Facebook’s open source applied reinforcement learning platform. *arXiv preprint arXiv:1811.00260*.
- Gulcehre, C., Paine, T. L., Shahriari, B., Denil, M., Hoffman, M., Soyer, H., Tanburn, R., Kapturowski, S., Rabinowitz, N., Williams, D., et al. (2020). Making efficient use of demonstrations to solve hard exploration problems. *8th International Conference on Learning Representations, ICLR 2020, Addis Adaba, France, April 26-1 May, 2020, Conference Track Proceedings*.



- Hessel, M., Modayil, J., van Hasselt, H., Schaul, T., Ostrovski, G., Dabney, W., Horgan, D., Piot, B., Azar, M. G., and Silver, D. (2018). Rainbow: Combining improvements in deep reinforcement learning. In McIlraith, S. A. and Weinberger, K. Q., editors, *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence, (AAAI-18), the 30th innovative Applications of Artificial Intelligence (IAAI-18), and the 8th AAAI Symposium on Educational Advances in Artificial Intelligence (EAAI-18), New Orleans, Louisiana, USA, February 2-7, 2018*, pages 3215–3222. AAAI Press.
- Hester, T., Vecerik, M., Pietquin, O., Lanctot, M., Schaul, T., Piot, B., Horgan, D., Quan, J., Sendonaris, A., Osband, I., et al. (2018). Deep Q-learning from demonstrations. In *AAAI Conference on Artificial Intelligence*.
- Hester, T., Vecerik, M., Pietquin, O., Lanctot, M., Schaul, T., Piot, B., Sendonaris, A., Dulac-Arnold, G., Osband, I., Agapiou, J., Leibo, J. Z., and Gruslys, A. (2017). Learning from demonstrations for real world reinforcement learning. *CoRR*, abs/1704.03732.
- Horgan, D., Quan, J., Budden, D., Barth-Maron, G., Hessel, M., van Hasselt, H., and Silver, D. (2018a). Distributed prioritized experience replay. In *International Conference on Learning Representations*.
- Horgan, D., Quan, J., Budden, D., Barth-Maron, G., Hessel, M., van Hasselt, H., and Silver, D. (2018b). Distributed prioritized experience replay.
- Jaderberg, M., Mnih, V., Czarnecki, W. M., Schaul, T., Leibo, J. Z., Silver, D., and Kavukcuoglu, K. (2017). Reinforcement learning with unsupervised auxiliary tasks. In *5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings*. OpenReview.net.
- Jaques, N., Lazaridou, A., Hughes, E., Gulcehre, C., Ortega, P., Strouse, D., Leibo, J. Z., and De Freitas, N. (2019). Social influence as intrinsic motivation for multi-agent deep reinforcement learning. In *International Conference on Machine Learning*, pages 3040–3049.
- Kapturowski, S., Ostrovski, G., Dabney, W., Quan, J., and Munos, R. (2019). Recurrent experience replay in distributed reinforcement learning. In *International Conference on Learning Representations*.
- Kim, B., Farahmand, A.-m., Pineau, J., and Precup, D. (2013). Learning from limited demonstrations. In *Advances in Neural Information Processing Systems*, pages 2859–2867.
- Kingma, D. and Ba, J. (2014). Adam optimizer. *arXiv preprint arXiv:1412.6980*, pages 1–15.
- Kulkarni, T. D., Narasimhan, K., Saeedi, A., and Tenenbaum, J. (2016). Hierarchical deep reinforcement learning: Integrating temporal abstraction and intrinsic motivation. In *Advances in Neural Information Processing Systems*, pages 3675–3683.
- Küttler, H., Nardelli, N., Lavril, T., Selvatici, M., Sivakumar, V., Rocktäschel, T., and Grefenstette, E. (2019). Torchbeast: A pytorch platform for distributed rl. *arXiv preprint arXiv:1910.03552*.
- Lagoudakis, M. G. and Parr, R. (2003). Least-squares policy iteration. *Journal of machine learning research*, 4(Dec):1107–1149.
- Lange, S., Gabel, T., and Riedmiller, M. (2012). Batch reinforcement learning. In *Reinforcement learning*, pages 45–73. Springer.
- Legg, S. and Hutter, M. (2007). Universal intelligence: A definition of machine intelligence. *Minds and Machines*, 17(4):391–444.
- Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., and Wierstra, D. (2016). Continuous control with deep reinforcement learning. In Bengio, Y. and LeCun, Y., editors, *4th International Conference on Learning Representations, ICLR 2016, San Juan, Puerto Rico, May 2-4, 2016, Conference Track Proceedings*.
- Lin, L.-J. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine learning*, 8(3-4):293–321.
- Michie, D. (1993). *Knowledge, Learning and Machine Intelligence*.

- Mnih, V., Badia, A. P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Silver, D., and Kavukcuoglu, K. (2016). Asynchronous methods for deep reinforcement learning. In Balcan, M. F. and Weinberger, K. Q., editors, *Proceedings of The 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pages 1928–1937, New York, New York, USA. PMLR.
- Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., et al. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540):529.
- Moritz, P., Nishihara, R., Wang, S., Tumanov, A., Liaw, R., Liang, E., Paul, W., Jordan, M. I., and Stoica, I. (2017). Ray: A distributed framework for emerging AI applications. *CoRR*, abs/1712.05889.
- Nair, A., McGrew, B., Andrychowicz, M., Zaremba, W., and Abbeel, P. (2018). Overcoming exploration in reinforcement learning with demonstrations. In *IEEE International Conference on Robotics and Automation*, pages 6292–6299.
- Nair, A., Srinivasan, P., Blackwell, S., Alcicek, C., Fearon, R., De Maria, A., Panneershelvam, V., Suleyman, M., Beattie, C., Petersen, S., et al. (2015). Massively parallel methods for deep reinforcement learning. *arXiv preprint arXiv:1507.04296*.
- OpenAI, Andrychowicz, M., Baker, B., Chociej, M., Józefowicz, R., McGrew, B., Pachocki, J. W., Pachocki, J., Petron, A., Plappert, M., Powell, G., Ray, A., Schneider, J., Sidor, S., Tobin, J., Welinder, P., Weng, L., and Zaremba, W. (2018). Learning dexterous in-hand manipulation. *arXiv preprint arXiv:1808.00177*.
- Osband, I., Aslanides, J., and Cassirer, A. (2018). Randomized prior functions for deep reinforcement learning. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, NIPS’18, page 8626–8638, Red Hook, NY, USA. Curran Associates Inc.
- Osband, I., Doron, Y., Hessel, M., Aslanides, J., Sezener, E., Saraiva, A., McKinney, K., Lattimore, T., Szepesvári, C., Singh, S., Roy, B. V., Sutton, R. S., Silver, D., and van Hasselt, H. (2020). Behaviour suite for reinforcement learning. In *International Conference on Representation Learning*.
- Pineau, J., Vincent-Lamarre, P., Sinha, K., Larivière, V., Beygelzimer, A., d’Alché-Buc, F., Fox, E. B., and Larochelle, H. (2020). Improving reproducibility in machine learning research (A report from the neurips 2019 reproducibility program). *CoRR*, abs/2003.12206.
- Piot, B., Geist, M., and Pietquin, O. (2014). Boosted Bellman residual minimization handling expert demonstrations. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 549–564. Springer.
- Pohlen, T., Piot, B., Hester, T., Azar, M. G., Horgan, D., Budden, D., Barth-Maroon, G., van Hasselt, H., Quan, J., Vecerík, M., Hessel, M., Munos, R., and Pietquin, O. (2018). Observe and look further: Achieving consistent performance on atari. *CoRR*, abs/1805.11593.
- Pomerleau, D. A. (1989). ALVINN: An autonomous land vehicle in a neural network. In *NIPS*, pages 305–313.
- Puterman, M. (1994). *Markov decision processes: discrete stochastic dynamic programming*. John Wiley & Sons.
- Russell, S. (2016). Rationality and intelligence: A brief update. In *Fundamental issues of artificial intelligence*, pages 7–28.
- Schaul, T., Quan, J., Antonoglou, I., and Silver, D. (2015). Prioritized experience replay.
- Schulman, J., Levine, S., Abbeel, P., Jordan, M., and Moritz, P. (2015). Trust region policy optimization. In *International conference on machine learning*, pages 1889–1897.
- Schulman, J., Moritz, P., Levine, S., Jordan, M., and Abbeel, P. (2016). High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations*.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.

- Sergio Guadarrama, Anoop Korattikara, O. R. P. C. E. H. S. F. K. W. E. G. N. W. C. H. V. V. E. B. (2018). TF-Agents: A library for reinforcement learning in tensorflow. <https://github.com/tensorflow/agents>. [Online; accessed 25-June-2019].
- Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., Van Den Driessche, G., Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., et al. (2016). Mastering the game of go with deep neural networks and tree search. *nature*, 529(7587):484.
- Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., Lillicrap, T., Simonyan, K., and Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362(6419):1140–1144.
- Silver, D., Lever, G., Heess, N., Degris, T., Wierstra, D., and Riedmiller, M. (2014). Deterministic policy gradient algorithms. In *Proceedings of the 32nd International Conference on Machine Learning, ICML, 22-24 June 2014, Beijing, China*.
- Song, Y., Wang, J., Lukasiewicz, T., Xu, Z., Xu, M., Ding, Z., and Wu, L. (2019). Arena: A general evaluation platform and building toolkit for multi-agent intelligence. *arXiv preprint arXiv:1905.08085*.
- Sutton, R. S. and Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.
- Sutton, R. S., McAllester, D. A., Singh, S. P., and Mansour, Y. (2000). Policy gradient methods for reinforcement learning with function approximation. In *Advances in neural information processing systems*, pages 1057–1063.
- Tassa, Y., Doron, Y., Muldal, A., Erez, T., Li, Y., de Las Casas, D., Budden, D., Abdolmaleki, A., Merel, J., Lefrancq, A., Lillicrap, T. P., and Riedmiller, M. A. (2018). Deepmind control suite. *CoRR*, abs/1801.00690.
- Todorov, E., Erez, T., and Tassa, Y. (2012). Mujoco: A physics engine for model-based control. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 5026–5033. IEEE.
- van Hasselt, H., Guez, A., and Silver, D. (2015). Deep reinforcement learning with double q-learning.
- van Hasselt, H., Hessel, M., and Aslanides, J. (2019). When to use parametric models in reinforcement learning? In Wallach, H. M., Larochelle, H., Beygelzimer, A., d’Alché-Buc, F., Fox, E. B., and Garnett, R., editors, *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, 8-14 December 2019, Vancouver, BC, Canada*, pages 14322–14333.
- Vecerik, M., Hester, T., Scholz, J., Wang, F., Pietquin, O., Piot, B., Heess, N., Rothörl, T., Lampe, T., and Riedmiller, M. A. (2017). Leveraging demonstrations for deep reinforcement learning on robotics problems with sparse rewards. *CoRR*, abs/1707.08817.
- Vezhnevets, A. S., Osindero, S., Schaul, T., Heess, N., Jaderberg, M., Silver, D., and Kavukcuoglu, K. (2017). Feudal networks for hierarchical reinforcement learning. In *Proceedings of the 34th International Conference on Machine Learning - Volume 70, ICML’17*, page 3540–3549. JMLR.org.
- Vinyals, O., Babuschkin, I., Czarnecki, W., Mathieu, M., Dudzik, A., Chung, J., Choi, D., Powell, R., Ewalds, T., Georgiev, P., Oh, J., Horgan, D., Kroiss, M., Danihelka, I., Huang, A., Sifre, L., Cai, T., Agapiou, J., Jaderberg, M., and Silver, D. (2019). Grandmaster level in starcraft ii using multi-agent reinforcement learning. *Nature*, 575.
- Wang, Z., Bapst, V., Heess, N., Mnih, V., Munos, R., Kavukcuoglu, K., and de Freitas, N. (2016). Sample efficient actor-critic with experience replay. *arXiv preprint arXiv:1611.01224*.
- Wang, Z., Schaul, T., Hessel, M., van Hasselt, H., Lanctot, M., and de Freitas, N. (2015). Dueling network architectures for deep reinforcement learning.
- Xu, Z., van Hasselt, H. P., and Silver, D. (2018). Meta-gradient reinforcement learning. In *Advances in Neural Information Processing Systems*, pages 2396–2407.
- Zhi, J., Wang, R., Clune, J., and Stanley, K. O. (2020). Fiber: A platform for efficient development and distributed training for reinforcement learning and population-based methods.

## A. Full benchmark results

### A.1. Control suite from features (all agents)




Figure 14 | Comparison of single process agent performance as measured by episode return with respect to actor steps.

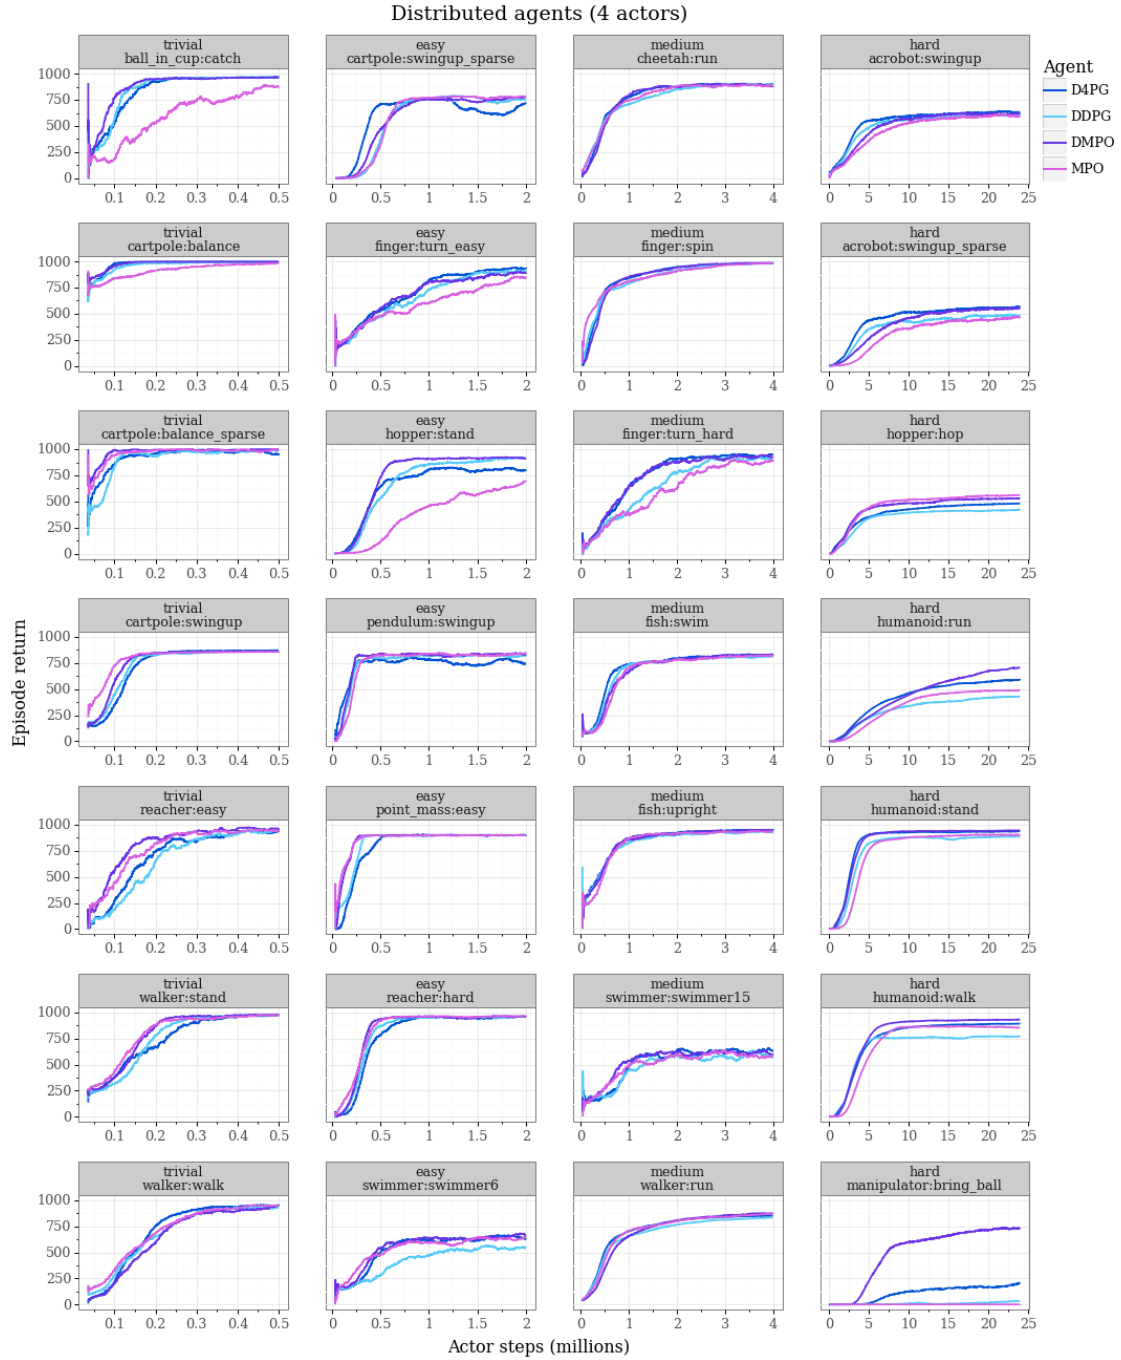


Figure 15 | Comparison of distributed agent performance as measured by episode return with respect to actor steps.

## A.2. Control suite from features (D4PG)

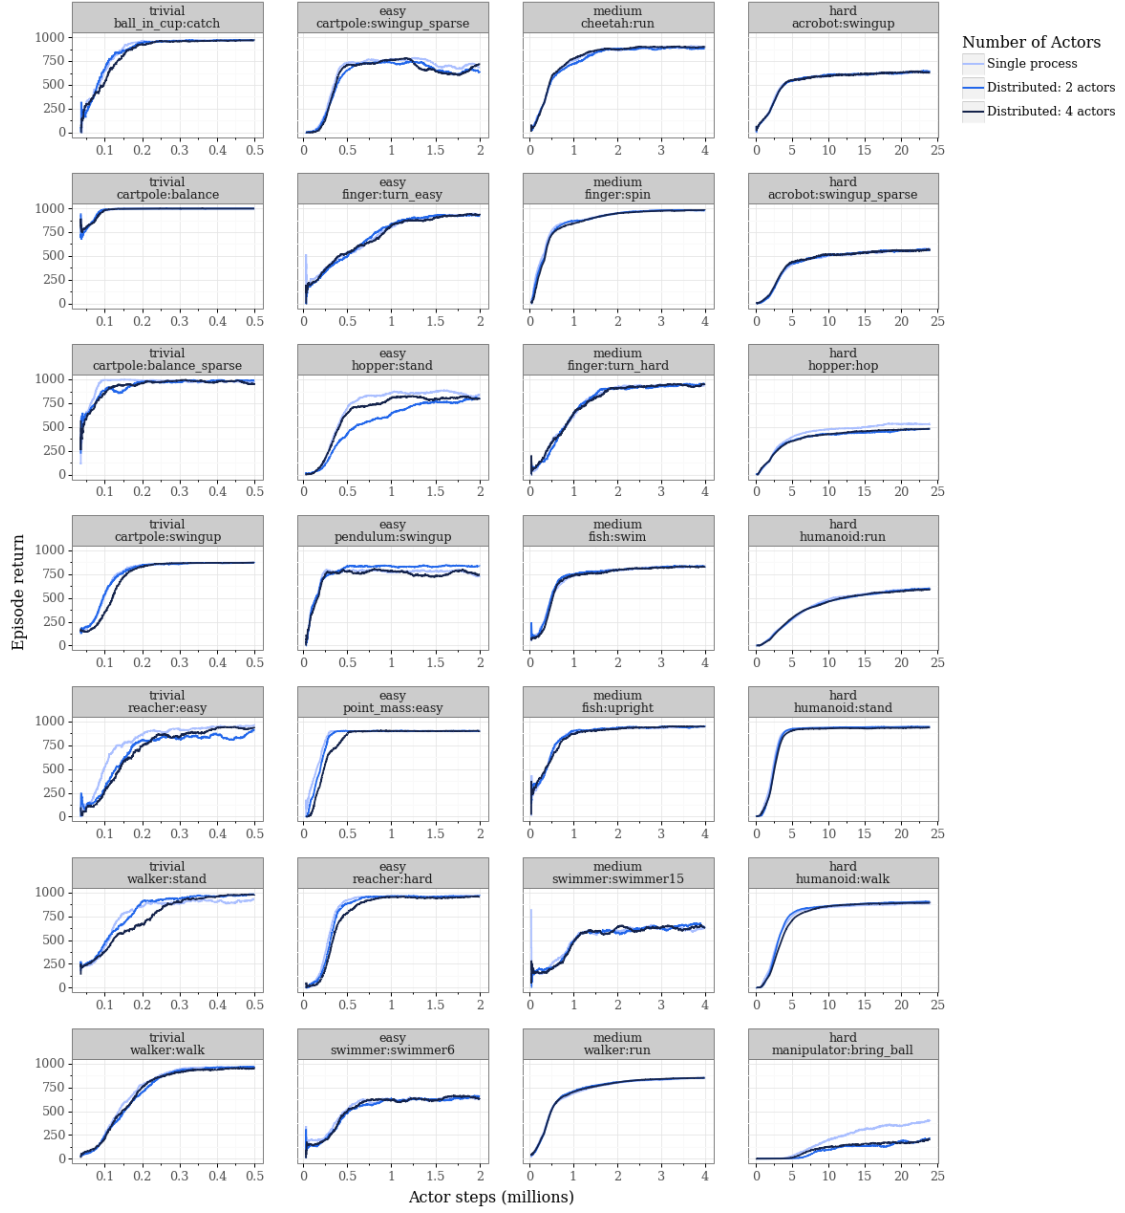


Figure 16 | Comparison of D4PG performance with different number of actors as measured by episode return with respect to actor steps.

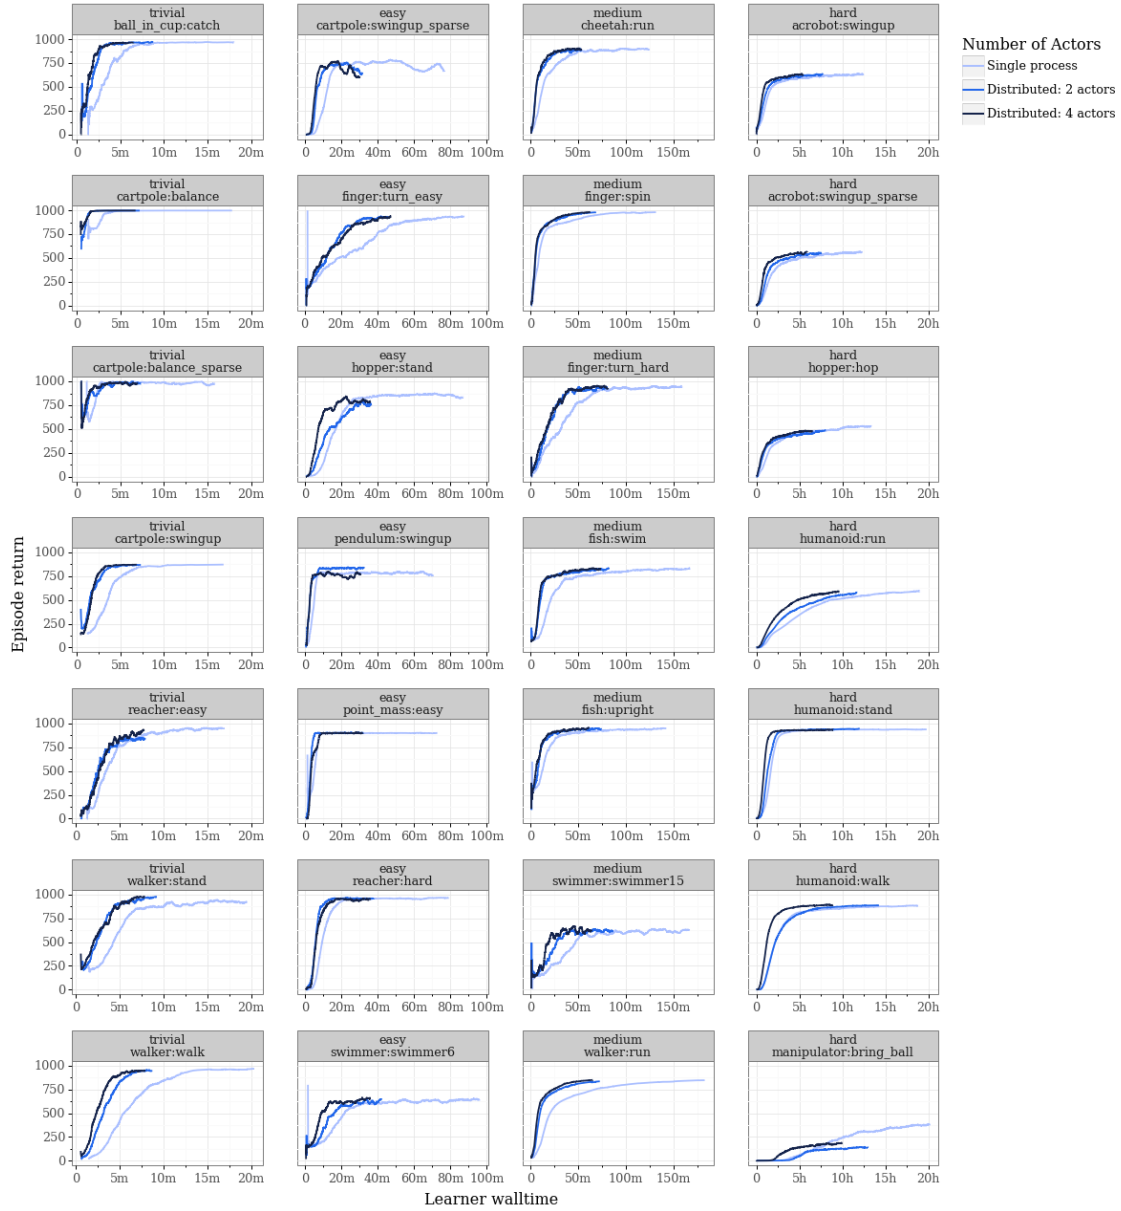


Figure 17 | Comparison of D4PG performance with different number of actors as measured by episode return with respect to learner walltime.



### A.3. Rate limiter sensitivity analysis

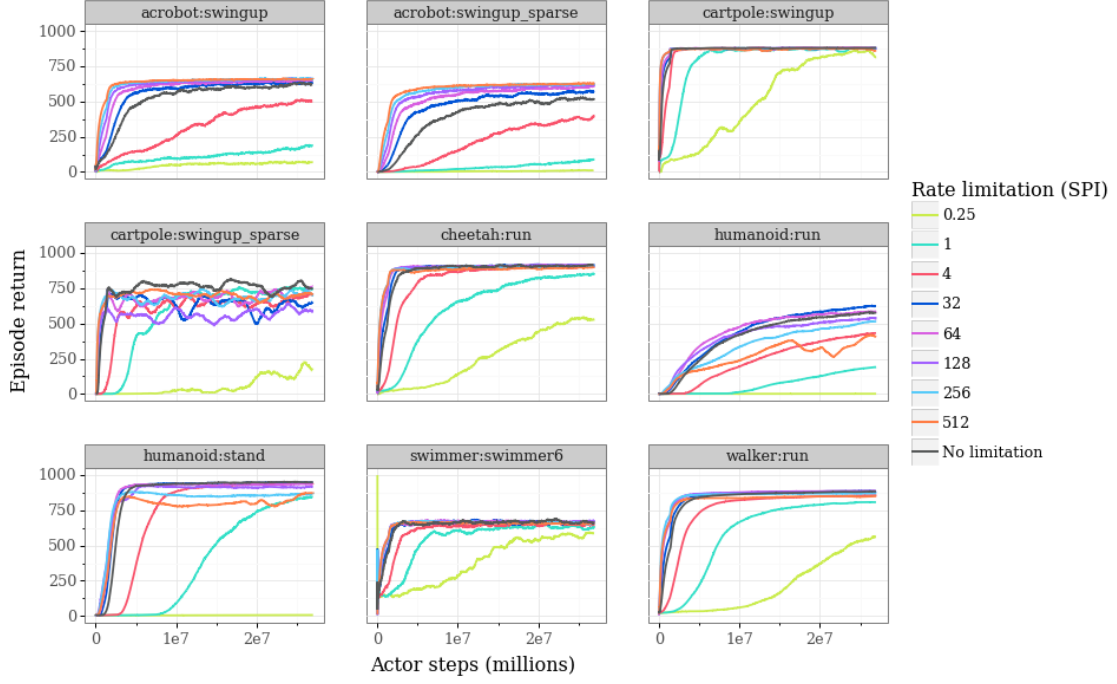


Figure 18 | Sensitivity analysis of the effect of rate limitation on the sample efficiency and performance of D4PG. All curves correspond to D4PG with 4 distributed actors. Episode returns are averaged over 10 seeds.

## A.4. Control suite from pixels (D4PG)




Figure 19 | Comparison of D4PG performance with pixel observations with different number of actors as measured by episode return with respect to actor steps.




Figure 20 | Comparison of D4PG performance with pixel observations with different number of actors as measured by episode return with respect to learner walltime.

## A.5. Atari (all agents)

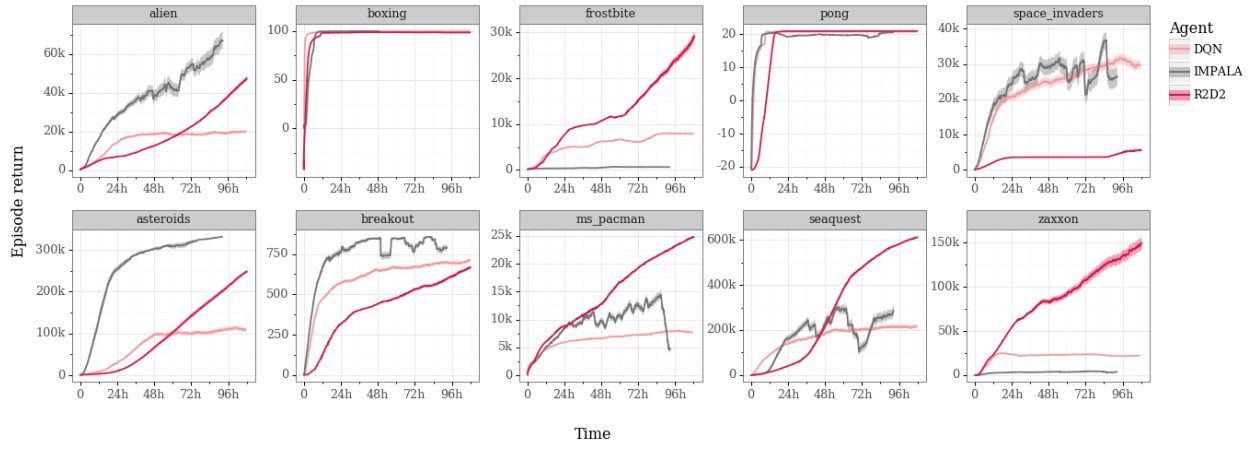


Figure 21 | Comparing performance of DQN, R2D2 and IMPALA on Atari tasks as measured by episode return with respect to training time.